

Sim-to-Sim Policy Transfer and Mismatch Analysis

Legged Locomotion across Three Robot Morphologies

Boonyaporn Preechasuth
VISTEC Internship Examination
bpbeam.boonyaporn@gmail.com

Abstract

Sim-to-sim policy transfer—deploying reinforcement learning policies trained in one physics simulator to another—provides a controlled testbed for understanding the sim-to-real gap without hardware confounds. This report presents a systematic analysis of transferring locomotion policies from Isaac Lab [6, 9] (NVIDIA PhysX) to MuJoCo [15] across three robot morphologies spanning the complexity spectrum: the **Unitree Go2** [16] quadruped (12 DOF, statically stable, uniform position control), the **Agility Cassie** [1] biped (10 actuated + 2 passive DOF, 4-bar linkage constraints), and the **Flamingo** two-wheel-legged robot (8 DOF, dynamically unstable, hybrid control with gear ratios and temporal stacking). I document the complete transfer pipeline for each platform, including the trial-and-error debugging process.

For **Go2** (Unitree Go2), four experiments show that moderate domain randomization (DR) is optimal (RMSE 0.257), while aggressive DR degrades coordination. The 39% RMSE increase is primarily due to actuator dynamics mismatch rather than contact physics. ActuatorNet and RMA effectively bridge this gap, with all variants achieving 100% task success.

For **Cassie** (Agility Cassie), the 4-bar linkage and passive ankle structure pose the main challenges. Fixing eight bugs in observation frames and PD scaling improved survival from 16 to 153 steps. Retraining with passive ankle constraints was necessary to achieve stable standing, proving that structural mismatches require architectural changes.

For **Flamingo**, eight compounding bugs, including wheel inversion and a 46% geometry mismatch, led to qualitative failure. Despite fixing the pipeline, the robot only balanced without locomotion. This highlights that dynamically unstable platforms amplify mismatches, and training/deployment geometry must match exactly.

Key findings: (1) Observation pipeline correctness dominates physics engine matching; (2) Dominant mismatch sources are morphology-dependent (actuator dynamics for Go2, kinematic structure for Cassie, control architecture for Flamingo); (3) DR, ActuatorNet, and RMA address orthogonal dimensions and should be combined; (4) Bugs are multiplicative, requiring critical mass before improvement; (5) Transfer difficulty scales super-linearly as each feature introduces interactions; (6) Geometry matching between training and deployment models is prerequisite for meaningful transfer.

Keywords: Sim-to-Sim Transfer · Domain Randomization · ActuatorNet · Rapid Motor Adaptation · Legged Locomotion · Isaac Lab · MuJoCo · Unitree Go2 · Cassie · Two-Wheel-Legged Robot

1 Introduction

Reinforcement learning has become the dominant paradigm for training agile locomotion controllers for legged robots [3, 7, 11]. The standard workflow trains policies in GPU-accelerated simulators such as Isaac Lab [6, 9] using algorithms like Proximal Policy Optimization (PPO) [12] across thousands of parallel environments, then transfers the learned policy to a target platform. When the target is physical hardware, this is *sim-to-real* transfer; when the target is a second simulator with different physics, this is *sim-to-sim* transfer.

Why study sim-to-sim transfer? It provides a controlled experimental framework for understanding transfer failures without hardware confounds. By deploying between two simulators with known—but different—physics engines (contact models, actuator dynamics, integrator schemes), I can isolate and quantify specific mismatch sources that are otherwise obscured by real-world noise, sensor latency, mechanical compliance, and model uncertainty. The insights from sim-to-sim transfer directly inform sim-to-real strategies: if a policy cannot transfer between two deterministic simulators, it will not transfer to stochastic hardware.

The role of domain randomization. Domain randomization (DR) [10, 14] addresses transfer brittleness by training policies on a *distribution* of environment parameters (friction, mass, actuator gains) rather than a single fixed configuration. The hypothesis is that a policy robust to training-time variation will generalize to deployment-time differences. However, DR’s effectiveness depends critically on *which* parameters are randomized and *how widely*—overly conservative DR undertrains robustness, while overly aggressive DR degrades task performance.

1.1 Robot Platforms and Transfer Complexity

This report studies sim-to-sim transfer across three robot morphologies that span the complexity spectrum:

1. **Unitree Go2** (Section 2): A 12-DOF quadruped with symmetric leg structure, uniform position control, and one-to-one joint mapping between simulators. Represents the “simplest” baseline transfer. *Key challenge*: Actuator dynamics mismatch (implicit vs. explicit PD). *Transfer quality*: 100% task success with 39% RMSE increase (3 bugs fixed).
2. **Agility Cassie** (Section 3): A 12-DOF biped with 4-bar linkage constraints that make 2 ankle joints passive in MuJoCo but actuated in Isaac Lab (12→10 action space mismatch). Represents a structurally mismatched transfer requiring environment redesign. *Key challenge*: Kinematic structure and constraint physics. *Transfer quality*: 1.70s standing with passive ankle retraining (8 bugs fixed).
3. **Flamingo** (Section 4): An 8-DOF two-wheel-legged robot with hybrid position-velocity control, mechanical gear ratios ($g = -1.5$), axis inversions, actuator delays (0–4 steps), and 3-frame temporal observation stacking. Dynamic instability (must actively balance on wheels) amplifies every observation and control mismatch. Represents the most complex transfer. *Key challenge*: Control mode coupling, actuator delay, and 46% geometry mismatch. *Transfer quality*: Basic balance only, no meaningful locomotion (8 bugs fixed + geometry blocker).


Complexity gradient. The three robots form an increasing gradient of transfer difficulty, enabling us to identify *which* types of mismatches are most damaging and *why*. Go2’s success establishes the baseline; Cassie’s partial success reveals the limits of pipeline debugging when structure differs; Flamingo’s failure demonstrates that dynamic instability and geometry mismatch create fundamental transfer barriers.

1.2 Methodology and Contributions

For each robot, I follow a four-stage methodology:

1. **Training Configuration:** Define environment, observation/action spaces, reward function, and control architecture in Isaac Lab.
2. **Debugging Journey:** Document the trial-and-error process of identifying and fixing transfer-breaking bugs (joint ordering, frame conventions, gain scaling, etc.).
3. **Controlled Experiments:** Test formal hypotheses with quantitative metrics across multiple scenarios (directional commands, terrain types, DR levels).
4. **Root Cause Analysis:** Decompose the sim-gap to identify the dominant mismatch source for each morphology.

The shared transfer pipeline is:



Key contributions: (1) Quantitative comparison of three transfer improvement methods (DR, ActuatorNet, RMA) on the same platform; (2) Systematic debugging documentation revealing that bugs are multiplicative, not additive; (3) Evidence that the dominant mismatch source is morphology-dependent, requiring tailored solutions; (4) Demonstration that geometry matching is prerequisite for dynamically unstable platforms; (5) Identification of transfer complexity scaling laws: each mechanical feature introduces super-linear interactions.

2 Part 1: Unitree Go2 Quadruped

2.1 Robot Overview

The Unitree Go2 [16] is a compact quadruped robot weighing approximately 15 kg with 12 degrees of freedom (3 per leg: hip abduction/adduction, thigh flexion, and calf flexion). Its symmetric four-legged structure and relatively low mass make it an ideal starting platform for sim-to-sim transfer studies. Quadrupeds have been widely used in reinforcement learning research for agile locomotion, ranging from trotting on flat ground to navigating rough terrain [7, 8, 11].

Table 1: Go2 Robot Specifications

Property	Value
Total mass	~15.0 kg (base body: 6.921 kg)
Degrees of freedom	12 (3 per leg: hip, thigh, calf)
Actuator type	DC motor with PD position control
Torque limit	23.7–45.4 N·m (joint-dependent)
Leg configuration	Symmetric: FL, FR, RL, RR
Standing height	~0.34–0.40 m (simulator-dependent)



Figure 1: Unitree Go2 quadruped robot in MuJoCo simulation.

Why Go2 first? The Go2’s 12-to-12 joint mapping between Isaac Lab and MuJoCo (no passive joints, no gear ratios, uniform actuator type) provides the cleanest baseline for validating the transfer pipeline before tackling more complex morphologies.

2.2 Transfer Pipeline

The Go2 transfer pipeline consists of four stages:

1. **Training** (Isaac Lab/PhysX): Train PPO policies with RSL-RL [11] in 4096 parallel environments on flat terrain. Three domain randomization levels are trained: Baseline, Moderate, and Aggressive.
2. **Export**: Convert trained PyTorch checkpoints to self-contained TorchScript (.pt) files via JIT tracing, removing the dependency on Isaac Lab for inference.
3. **Deployment** (MuJoCo): Load the TorchScript policy in a standalone MuJoCo simulation. At each control step: read MuJoCo state → reorder joints to Isaac Lab convention → construct 48-dim observation → run policy forward pass → compute PD torques → apply control swap → step physics.
4. **Evaluation**: Run headless evaluation across 8 directional scenarios (5 episodes × 500 steps each) and compute tracking metrics.

Table 2: Go2 Policies Trained and Evaluated

Policy	DR Level	Iterations	Experiment Name
baseline.pt	Minimal (friction only)	5,000	go2_sim2sim
moderate.pt	Moderate (mass, gains, push)	8,000	go2_sim2sim_moderate
aggressive.pt	Aggressive (wide ranges)	10,000	go2_sim2sim_aggressive

2.3 Training Journey: From Complete Failure to 100% Success

The path to successful sim-to-sim transfer was far from straightforward. Initial transfer attempts resulted in **complete failure**—the robot collapsed immediately upon deployment in MuJoCo and could not even maintain a standing pose with zero velocity commands.

2.3.1 Phase 1: Initial Failure (All Attempts Failed)

The first deployment attempts produced a robot that collapsed within the first 0.1 seconds in MuJoCo. Extensive debugging attempts were made, none of which resolved the issue:

- **Height adjustments:** Varied starting height from 0.35m to 0.42m—no improvement.
- **Gain increases:** Tried $K_p = 50, 100, 200$ —worsened oscillation.
- **DCMotor saturation:** Implemented velocity-dependent torque saturation matching Isaac Lab’s model—no improvement.

- **Zero command test:** Even commanding the robot to stand still ($v_x = v_y = \omega_z = 0$) resulted in immediate collapse.

The zero-command test was the critical diagnostic: if the robot cannot stand still, the problem is not in command tracking but in the *fundamental deployment pipeline*—observation construction, joint mapping, or actuator configuration.

2.3.2 Phase 2: Root Cause Identification

Systematic investigation revealed three independent, simultaneously-occurring bugs:

Table 3: Root Causes of Go2 Transfer Failure

Bug	Wrong Value	Correct Value	Impact
PD stiffness (K_p)	25.0	20.0	25% gain mismatch $\rightarrow$ joint overshoot $\rightarrow$ unstable oscillation
Control signal swap	Not applied	FR $\leftrightarrow$ FL, RR $\leftrightarrow$ RL	Torques sent to contralateral legs $\rightarrow$ immediate collapse
MuJoCo XML source	mujoco_menagerie	unitree_mujoco	Extra damping=2.0, frictionloss=0.2, lower torque limits

Bug 1: PD Gain Mismatch ($K_p = 25$ instead of 20). The Isaac Lab environment configuration specifies `DCMotorCfg` with `stiffness=20.0` and `damping=0.5`. The initial deployment script used $K_p = 25.0$, sourced from an older configuration file. This 25% gain mismatch caused joints to overshoot their target positions, creating oscillations that compounded across the 12 joints.

Bug 2: Actuator Ordering Mismatch (Missing Control Swap). MuJoCo maintains *two separate orderings*: joint positions in `qpos` follow the order FL, FR, RL, RR (by leg), while actuator control signals in `ctrl` follow FR, FL, RR, RL (as defined in the XML `<actuator>` section). Without the control swap, torques computed for the front-left leg were applied to the front-right leg and vice versa, producing physically nonsensical forces:

```
# MuJoCo actuator order: FR, FL, RR, RL
# Isaac Lab policy output order: FL, FR, RL, RR
d.ctrl[0:3] = tau[3:6]      # FR actuators <- FL torques
d.ctrl[3:6] = tau[0:3]      # FL actuators <- FR torques (swap!)
d.ctrl[6:9] = tau[9:12]     # RR actuators <- RL torques
d.ctrl[9:12] = tau[6:9]     # RL actuators <- RR torques (swap!)
```

Listing 1: Correct control signal swap for Go2 MuJoCo deployment

Bug 3: Wrong MuJoCo XML Model. The `mujoco_menagerie` version of the Go2 model contained additional dynamics parameters not present in the training simulator:

- Joint damping=2.0 (Isaac Lab default: 0.0, with friction handled by the actuator model)
- Joint frictionloss=0.2 (not modeled in Isaac Lab)
- Lower `ctrlrange` limits (± 23.7 N·m vs. ± 30.0 N·m in Isaac Lab)

Switching to the official `unitree_mujoco` XML (`scene_flat.xml`) with `damping=0.1` and `armature=0.01` provided a much closer physics match.

2.3.3 Phase 3: Successful Transfer

After applying all three fixes simultaneously, all three DR policies achieved **100% success rate** across all test scenarios. The key lesson: sim-to-sim transfer bugs are often *multiplicative*—each individual bug may not be immediately diagnosable because the combined effect masks the individual contributions. Fixing bugs one at a time may not show improvement until *all* are resolved.

Standing height validation served as the final sanity check. The same default joint angles produced different standing heights across simulators:

- Isaac Lab (PhysX): ~ 0.40 m
- MuJoCo (wrong XML): ~ 0.27 – 0.29 m
- MuJoCo (correct XML): ~ 0.34 m

The residual 0.06 m difference between Isaac Lab and the correct MuJoCo XML reflects inherent differences in collision geometry, contact solver stiffness, and joint dynamics between PhysX and MuJoCo.

2.4 Model Architecture and Training Configuration

2.4.1 Neural Network Architecture

All three Go2 policies share identical network architecture, trained with RSL-RL’s PPO implementation [11, 12]:

Table 4: Go2 Policy Network Architecture

Component	Specification
Actor	$48 \rightarrow [512, 256, 128] \rightarrow 12$ (ELU activations)
Critic	$48 \rightarrow [512, 256, 128] \rightarrow 1$ (ELU activations)
Init noise std	1.0
Input dimension	48 (observation vector)
Output dimension	12 (joint position offsets, scaled by 0.25)
File size	759 KB (TorchScript)

2.4.2 PPO Hyperparameters

Table 5: PPO Training Hyperparameters (Go2)

Parameter	Value	Parameter	Value
Algorithm	PPO	Clip parameter	0.2
Steps per env	24	Entropy coefficient	0.01
Mini-batches	4	Discount (γ)	0.99
Learning epochs	5	GAE λ	0.95
Learning rate	$1.0\text{e-}3$ (adaptive)	Desired KL	0.01
Num environments	4,096	Max grad norm	1.0
Empirical normalization	No		

2.4.3 Observation Space (48 dimensions)

The policy receives a 48-dimensional observation vector constructed identically in both simulators:

Table 6: Go2 Observation Space (48 dimensions)

Index	Name	Dim	Scale	Description
0–2	base_lin_vel	3	$\times 2.0$	Base linear velocity (body frame)
3–5	base_ang_vel	3	$\times 0.25$	Base angular velocity (body frame)
6–8	projected_gravity	3	$\times 1.0$	Gravity projected into body frame
9–11	velocity_commands	3	$\times 1.0$	Commanded $[v_x, v_y, \omega_z]$
12–23	joint_pos_rel	12	$\times 1.0$	Joint angles relative to default
24–35	joint_vel	12	$\times 0.05$	Joint velocities
36–47	last_action	12	$\times 1.0$	Previous policy output

Action space: The policy outputs 12-dimensional joint position *offsets*. These are scaled by $\alpha = 0.25$ and added to default joint angles to form PD position targets:

$$q_{\text{target}} = q_{\text{default}} + \alpha \cdot a_t, \quad \tau = K_p(q_{\text{target}} - q_t) - K_d\dot{q}_t \quad (1)$$

2.4.4 Reward Function

Table 7: Go2 Reward Function

Term	Weight	Purpose
track_lin_vel_xy_exp	+1.0	Primary: exponential linear velocity tracking ($\sigma = 0.5$)
track_ang_vel_z_exp	+0.5	Primary: exponential angular velocity tracking ($\sigma = 0.5$)
feet_air_time	+0.125	Encourage dynamic gait (threshold 0.5s)
lin_vel_z_l2	-2.0	Penalize vertical bouncing
ang_vel_xy_l2	-0.05	Penalize body roll/pitch rate
dof_torques_l2	-0.0002	Torque minimization
dof_vel_l2	-0.0001	Joint velocity regularization
dof_acc_l2	-2.5×10^{-7}	Joint acceleration smoothness
action_rate_l2	-0.01	Action smoothness
dof_pos_limits	-10.0	Joint limit violation avoidance
undesired_contacts	-1.0	Penalize thigh/calf ground contact

2.5 Simulator Environment Comparison

2.5.1 Isaac Lab (Training Simulator)

- **Physics engine:** NVIDIA PhysX (GPU-accelerated)
- **Timestep:** 0.005 s (200 Hz), decimation 4 $\rightarrow$ policy at 50 Hz
- **Actuator model:** DCMotorCfg (implicit PD). $K_p = 20.0$, $K_d = 0.5$, saturation 30.0 N·m. The implicit model means torques are computed *within* the physics solver, allowing simultaneous resolution of contact and actuator forces.
- **Episode length:** 20.0 s (1,000 steps at 50 Hz)
- **Terrain:** Flat plane (static friction = 1.0, dynamic friction = 1.0)
- **Initial height:** 0.42 m
- **Contact solver:** PhysX GPU TGS (iterative)

2.5.2 MuJoCo (Deployment Simulator)

- **Physics engine:** MuJoCo [15] (CPU)

- **Model file:** `unitree_mujoco/unitree_robots/go2/scene_flat.xml`
- **Timestep:** 0.005 s (200 Hz), decimation 4 $\rightarrow$ policy at 50 Hz
- **Actuator model:** Explicit PD implemented in Python. Torques computed *before* the physics step, introducing a one-step delay between error measurement and force application.
- **Contact solver:** MuJoCo convex Newton method with soft contacts
- **Integrator:** Semi-implicit Euler

2.5.3 Key Physics Differences

Table 8: Physics Engine Differences: Isaac Lab vs. MuJoCo (Go2)

Property	Isaac Lab (PhysX)	MuJoCo
Contact solver	GPU TGS (iterative)	Convex Newton (direct)
Friction model	Pyramidal approximation	Elliptic cone
Integration	Implicit Euler	Semi-implicit Euler
PD actuator	Implicit (solver-integrated)	Explicit (pre-step torque)
Joint damping	0.0 (actuator handles it)	0.1 (from XML)
Joint armature	Not explicit	0.01
Friction loss	Not modeled	0.0 (unitree_mujoco)
Torque limits	30.0 N·m (uniform)	23.7/45.4 N·m (variable)
Standing height	~ 0.40 m	~ 0.34 m

2.6 Joint Ordering and Observation Matching

2.6.1 Joint Ordering Mismatch

Isaac Lab and MuJoCo use fundamentally different joint indexing conventions. Isaac Lab groups joints **by type** (all hips, then all thighs, then all calves), while MuJoCo groups joints **by leg** (all joints of FL, then FR, then RL, then RR):

Table 9: Joint Ordering: Isaac Lab vs. MuJoCo

Index	Isaac Lab	MuJoCo
0	FL_hip	FL_hip
1	FR_hip	FL_thigh
2	RL_hip	FL_calf
3	RR_hip	FR_hip
4	FL_thigh	FR_thigh
5	FR_thigh	FR_calf
6	RL_thigh	RL_hip
7	RR_thigh	RL_thigh
8	FL_calf	RL_calf
9	FR_calf	RR_hip
10	RL_calf	RR_thigh
11	RR_calf	RR_calf

The conversion is handled by two index arrays:

```
MUJOCO_TO_ISAACLAB = [0, 3, 6, 9, 1, 4, 7, 10, 2, 5, 8, 11]
ISACLAB_TO_MUJOCO = [0, 4, 8, 1, 5, 9, 2, 6, 10, 3, 7, 11]
```


2.6.2 Observation Construction in MuJoCo

For the policy to produce correct outputs, the 48-dimensional observation must be constructed *identically* to how Isaac Lab constructs it internally. The MuJoCo deployment script performs these steps at each 50 Hz control cycle:

1. **Read raw state:** Extract quaternion orientation (`qpos[3:7]`), world-frame linear velocity (`qvel[0:3]`), world-frame angular velocity (`qvel[3:6]`), and joint states (`qpos[7:]`, `qvel[6:]`).
2. **Frame transformation:** Rotate world-frame linear velocity to body frame using the `quat_rotate_inverse` function. For MuJoCo free joints, angular velocity in `qvel[3:6]` is in world frame, so it also requires rotation.
3. **Joint reordering:** Apply `MUJOCO_TO_ISAACLAB` to convert joint positions and velocities from MuJoCo’s by-leg ordering to Isaac Lab’s by-type ordering.
4. **Scaling:** Apply the exact scaling factors used during training: $\times 2.0$ for linear velocity, $\times 0.25$ for angular velocity, $\times 0.05$ for joint velocities.
5. **Gravity projection:** Compute the gravity vector in the body frame using the quaternion orientation.

The `quat_rotate_inverse` function computes $R(q)^{-1} \cdot v$, where $R(q)$ is the rotation matrix corresponding to quaternion q :

$$v_{\text{body}} = v_{\text{world}} - 2w \cdot (q_{\text{xyz}} \times v_{\text{world}}) + 2(q_{\text{xyz}} \times (q_{\text{xyz}} \times v_{\text{world}})) \quad (2)$$

where $q = [w, x, y, z]$ and $q_{\text{xyz}} = [x, y, z]$. **The sign before w is critical:** using $+w$ gives the *forward* rotation (incorrect), while $-w$ gives the *inverse* (correct).

2.7 Domain Randomization Configurations

Three levels of domain randomization were designed to test the robustness-precision trade-off. The **only** variable across the three training runs is the DR configuration; all other parameters (reward weights, network architecture, observation structure) remain identical.

Table 10: Domain Randomization Configurations (Go2)

Parameter	Baseline	Moderate	Aggressive
Friction (static & dynamic)	[0.8, 1.2]	[0.5, 1.5]	[0.3, 1.8]
Base mass offset	None	$[-0.5, +1.0]$ kg	$[-1.0, +2.0]$ kg
K_p scaling	Fixed (1.0 $\times$)	$[0.8, 1.2] \times$	$[0.6, 1.4] \times$
K_d scaling	Fixed (1.0 $\times$)	$[0.8, 1.2] \times$	$[0.6, 1.4] \times$
External push velocity	None	± 0.5 m/s	± 0.7 m/s
Push interval	N/A	10–15 s	8–12 s
Joint reset range	$[0.5, 1.5] \times$	$[0.5, 1.5] \times$	$[0.3, 1.7] \times$
Training iterations	5,000	8,000	10,000
Total env interactions	$\sim 4.9\text{M}$	$\sim 7.9\text{M}$	$\sim 9.8\text{M}$

Moderate and Aggressive DR use proportionally more training iterations to compensate for the harder optimization landscape introduced by wider parameter ranges.

2.8 Flat Terrain Experiments: Baseline Transfer and Enhancement Strategies

The following four experiments evaluate policy transfer on **flat terrain** in MuJoCo, matching the training environment in Isaac Lab. These experiments establish the baseline transfer quality and

explore three enhancement strategies: Domain Randomization (DR), learned actuator models (ActuatorNet), and online adaptation (RMA).

Common Setup for Experiments 1A–1D:

- **Terrain:** Flat ground (`scene_flat.xml`) in both Isaac Lab (training) and MuJoCo (deployment)
- **Evaluation:** 8 directional scenarios (Forward, Backward, Left, Right, Turn_Left, Turn_Right, Diagonal, Circle).
- **Metrics:** Success rate, velocity tracking RMSE (v_x, v_y, ω_z), action smoothness, torque magnitude
- **Episodes:** 5 trials per scenario, 500 steps (10 seconds) per episode

2.9 Experiment 1A: Effect of Domain Randomization on Transfer Quality

2.9.1 Introduction and Motivation

Domain randomization during training is the primary defense against the sim-to-sim (and sim-to-real) gap [10, 14]. By exposing the policy to a distribution of environment parameters rather than a fixed configuration, DR forces the policy to learn robust behaviors that generalize beyond the training dynamics. However, wider randomization ranges come with a cost: the policy must be conservative to handle extreme parameter combinations, potentially sacrificing precision on the nominal (most likely) configuration.

This experiment compares three levels of DR for the Go2 quadruped on flat terrain, all deployed in MuJoCo after training in Isaac Lab.

2.9.2 Hypotheses

Hypothesis 1 (Distributional Robustness). *Policies trained with higher domain randomization will achieve higher survival rates when transferred to MuJoCo, as the broader training distribution better covers the MuJoCo dynamics mismatch.*

Hypothesis 2 (Velocity Tracking Trade-off). *Higher DR will degrade velocity tracking precision (higher RMSE), since the policy must adopt conservative behaviors to handle a wider range of dynamics—a robustness-precision trade-off.*

Hypothesis 3 (Locomotion Quality Shift). *Aggressive DR policies will exhibit less smooth actions and potentially different torque profiles, as the policy learns more reactive, corrective behavior to handle perturbations during training.*

2.9.3 Experiment Setup

Evaluation protocol: Each of the three policies (Baseline, Moderate, Aggressive) was evaluated using headless MuJoCo simulation (`eval_mujoco.py`) that exactly replicates the deployment pipeline:

- Joint mapping (Isaac Lab $\leftrightarrow$ MuJoCo) with control swap
- Identical 48-dim observation construction with matching scaling factors
- PD control: $K_p = 20.0$, $K_d = 0.5$, applied at 50 Hz
- MuJoCo model: `unitree_mujoco/unitree_robots/go2/scene_flat.xml`

Test scenarios (8 directional commands):

Table 11: Evaluation Scenarios (Go2)

Scenario	Command $[v_x, v_y, \omega_z]$	Axis	Description
Forward	[0.5, 0.0, 0.0]	v_x	Standard forward walking
Backward	[-0.5, 0.0, 0.0]	v_x	Reverse walking
Left	[0.0, 0.5, 0.0]	v_y	Lateral left
Right	[0.0, -0.5, 0.0]	v_y	Lateral right
Turn_Left	[0.0, 0.0, 0.5]	ω_z	In-place rotation
Turn_Right	[0.0, 0.0, -0.5]	ω_z	In-place rotation
Diagonal	[0.5, 0.5, 0.0]	$v_x + v_y$	Multi-axis: forward-left
Circle	[0.5, 0.0, 0.3]	$v_x + \omega_z$	Multi-axis: forward with turning

Per scenario: 5 episodes $\times$ 500 steps (10 s at 50 Hz). Success is defined as surviving $\geq 80\%$ of max steps (≥ 400 steps).

Metrics:

- **Success rate:** Percentage of episodes surviving ≥ 400 steps
- **RMSE** (v_x, v_y, ω_z , overall): Root mean square error between achieved and commanded velocity
- **Mean body height:** Average CoM height during locomotion
- **Action smoothness:** Mean L2 norm of consecutive action differences (lower = smoother)
- **Mean torque magnitude:** Mean L2 norm of applied torques

2.9.4 Results

Survival and Success Rate All three DR levels achieved **100% success rate** (40/40 episodes) across all 8 scenarios, with every episode reaching the full 500 steps. No falls were observed for any policy.

Table 12: Success Rate Summary (Experiment 1A)

DR Level	Overall Success	Mean Survival Steps	Falls
Baseline	100% (40/40)	500.0	0
Moderate	100% (40/40)	500.0	0
Aggressive	100% (40/40)	500.0	0

Table 13: Overall RMSE by DR Level (averaged across 8 scenarios)

DR Level	Overall RMSE	Mean RMSE v_x	Mean RMSE v_y	Mean RMSE ω_z
Baseline	0.273	0.223	0.168	0.201
Moderate	0.257	0.206	0.154	0.197
Aggressive	0.271	0.200	0.143	0.274

Table 14: Per-Scenario Overall RMSE (bold = best for that scenario)


Scenario	Baseline	Moderate	Aggressive	Best
Forward	0.238	0.229	0.240	Moderate
Backward	0.269	0.247	0.222	Aggressive
Left	0.249	0.223	0.215	Aggressive
Right	0.231	0.220	0.201	Aggressive
Turn_Left	0.287	0.284	0.271	Aggressive
Turn_Right	0.288	0.286	0.285	Aggressive
Diagonal	0.345	0.302	0.453	Moderate
Circle	0.277	0.267	0.279	Moderate
Average	0.273	0.257	0.271	Moderate

Velocity Tracking Accuracy (RMSE)


Table 15: Locomotion Quality Metrics (Go2 DR Comparison)

DR Level	Action Smoothness*	Mean Torque (N·m)	Mean Height (m)
Baseline	0.724 (smoothest)	15.2	0.343
Moderate	0.759	14.2	0.357
Aggressive	0.963 (most reactive)	13.8	0.365


* Lower = smoother. Defined as mean $\|a_t - a_{t-1}\|_2$.

Go2 DR Comparison — Overall

(a) Overall comparison



(b) Per-scenario and per-component RMSE



(c) Velocity tracking timeseries

Figure 2: Domain Randomization comparison across three DR levels (Baseline, Moderate, Aggressive). **(a)** Overall: all achieve 100% success; Moderate best RMSE (0.257). **(b)** Per-scenario breakdown: Aggressive fails on Diagonal multi-axis scenario. **(c)** Timeseries: Diagonal reveals yaw instability (ω_z oscillation ~ 1.0 rad/s) in Aggressive.

Locomotion Quality

2.9.5 Analysis and Discussion

Observation 1: Flat terrain does not discriminate survival. All three policies achieved 100% survival, meaning the **Distributional Robustness** hypothesis *cannot be evaluated* on this test case. The flat terrain MuJoCo transfer is a relatively benign scenario—the physics gap between PhysX and MuJoCo at moderate speeds (0.5 m/s) on flat ground is small enough that even minimal DR covers it. This hypothesis requires testing on rough terrain, with external perturbations, or at higher speeds to be properly validated.

Observation 2: Moderate DR is the optimal regularizer. Contrary to the simple prediction of Hypothesis 2 (monotonic precision degradation with more DR), Moderate DR actually *improved* overall tracking compared to Baseline (RMSE 0.257 vs. 0.273). This suggests

that mild randomization acts as a regularizer that helps the policy generalize across the simulator gap, analogous to how dropout regularization can improve test-set accuracy despite “degrading” training-set performance. The improvement is consistent across most scenarios (Moderate wins in 6/8).

Observation 3: Aggressive DR breaks multi-axis coordination. Aggressive DR improved single-axis tracking (v_x : 0.200 vs. 0.223 Baseline; v_y : 0.143 vs. 0.168 Baseline) but *catastrophically degraded* the Diagonal scenario (ω_z RMSE = 0.645, compared to 0.093 for Moderate—a $6.9\times$ increase). Time-series analysis reveals that the Aggressive policy exhibits large yaw rate spikes (~ 1.0 rad/s) during diagonal locomotion, suggesting that the strong perturbation-response mechanisms learned during training activate inappropriately during steady multi-axis motion.

This partially supports the **robustness-precision trade-off** hypothesis: the trade-off is real, but manifests specifically in *multi-axis coordination* rather than uniformly across all scenarios.

Observation 4: DR changes the locomotion strategy. The **Locomotion Quality Shift** hypothesis is **supported**. The Aggressive policy exhibits:

- 33% higher action rate (0.963 vs. 0.724), consistent with more reactive, corrective control learned to handle unpredictable perturbations
- 9% lower average torque (13.8 vs. 15.2 N·m), suggesting a more cautious, energy-efficient gait
- 6.4% higher body height (0.365 m vs. 0.343 m), indicating a taller stance adopted as a stability margin against the wider range of dynamics experienced during training

These changes represent a fundamental shift in *locomotion strategy*: the Aggressive policy trades smoothness for reactivity and adopts a taller, lower-energy posture.

2.9.6 Conclusion (Experiment 1A)

1. **Robustness:** Not differentiable—all policies survive flat terrain. Requires harder test conditions.
2. **Precision trade-off:** Partially supported. Moderate DR improves precision (regularization effect). Aggressive DR improves single-axis but breaks multi-axis coordination.
3. **Quality shift:** Supported. Higher DR produces more reactive actions, lower torque, and taller stance.
4. **Practical recommendation:** Moderate DR (RMSE 0.257) offers the best balance for flat terrain sim-to-sim deployment.

2.10 Experiment 1B: Sim-to-Sim Transfer Fidelity

2.10.1 Introduction and Motivation

Experiment 1A compared DR levels within MuJoCo only. A natural follow-up question is: *how faithfully does MuJoCo reproduce the policy’s behavior in its training simulator (Isaac Lab)?* If the same policy produces similar velocity profiles in both simulators, the sim-gap is small and the transfer is high-fidelity. If the profiles diverge significantly, the gap highlights specific dynamics mismatches that DR must compensate for.

This experiment runs the **same Baseline policy** in both Isaac Lab (PhysX) and MuJoCo under identical command scenarios, then directly overlays the resulting velocity time-series.

2.10.2 Hypotheses

Hypothesis 4 (Transfer Fidelity). *The Baseline policy will track velocity commands more accurately in Isaac Lab than in MuJoCo, since the policy was optimized for PhysX dynamics.*

Hypothesis 5 (Oscillation Amplification). *MuJoCo’s different integrator and contact model will amplify oscillations in the velocity signal compared to Isaac Lab’s smoother PhysX dynamics, due to the explicit PD controller introducing a one-step delay.*

Hypothesis 6 (Contact Model Divergence). *The largest gap will appear in angular velocity (ω_z) due to torsional friction differences between PhysX’s pyramidal approximation and MuJoCo’s elliptic friction cone.*

2.10.3 Experiment Setup

- **Policy:** Baseline (minimal DR, 5,000 iterations)
- **Isaac Lab evaluation:** `eval_isaaclab.py`—headless, 8 scenarios, 5 episodes $\times$ 500 steps
- **MuJoCo evaluation:** `eval_mujoco.py`—headless, same 8 scenarios, 5 episodes $\times$ 500 steps
- **Comparison method:** Best episode per scenario overlaid on the same axes

Both evaluators share identical observation construction (48-dim, same scaling), the same exported `policy.pt` checkpoint, fixed velocity commands per scenario, height-based termination (< 0.15 m), and 50 Hz recording rate. The **only difference** is the underlying physics engine.

2.10.4 Results

Success Rate Both simulators achieve **100% success rate** across all 8 scenarios.

Velocity Tracking RMSE MuJoCo exhibits higher tracking error with an average RMSE of 0.273, a 39% increase over the 0.197 recorded in Isaac Lab.

Table 16: Sim-to-Sim RMSE: Isaac Lab vs. MuJoCo (Baseline Policy)

Scenario	Isaac Lab	MuJoCo	Ratio (MJ/IL)
Forward	0.142	0.238	1.68 $\times$
Backward	0.167	0.269	1.61 $\times$
Left	0.164	0.249	1.52 $\times$
Right	0.190	0.231	1.22 $\times$
Turn_Left	0.278	0.287	1.03 $\times$
Turn_Right	0.299	0.288	0.96 $\times$
Diagonal	0.109	0.345	3.17 $\times$
Circle	0.227	0.277	1.22 $\times$
Average	0.197	0.273	1.39$\times$

Table 17: Per-Component RMSE Comparison

Component	Isaac Lab	MuJoCo	Ratio (MJ/IL)
RMSE v_x (mean)	0.076	0.223	2.93 $\times$
RMSE v_y (mean)	0.080	0.168	2.10 $\times$
RMSE ω_z (mean)	0.305	0.201	0.66 $\times$

Locomotion Quality Isaac Lab demonstrates superior quality, with 27% smoother actions, 17% lower torque consumption, and a 1.6 cm higher ground clearance.

Table 18: Locomotion Quality: Isaac Lab vs. MuJoCo

Metric	Isaac Lab	MuJoCo	Observation
Action smoothness	0.531	0.724	Isaac Lab 27% smoother
Mean torque (N·m)	12.7	15.2	Isaac Lab 17% less torque
Mean height (m)	0.359	0.343	Isaac Lab 1.6 cm taller




Figure 3: Sim-to-Sim velocity tracking overlay for baseline DR policy. **Red/solid:** Isaac Lab (training simulator). **Blue/dashed:** MuJoCo (deployment simulator). Forward (top), Left (middle), Diagonal (bottom). The dramatically larger oscillation amplitude in MuJoCo’s linear velocity signals (39% RMSE increase) confirms Hypothesis 5. Despite the sim-gap, both achieve 100% survival.

2.10.5 Analysis and Discussion

Transfer Fidelity—Supported for linear velocity, reversed for angular. Linear velocity tracking (v_x , v_y) is 2–3 $\times$ better in Isaac Lab (RMSE 0.076/0.080 vs. 0.223/0.168), confirming that the policy is optimized for PhysX dynamics. However, *angular velocity* (ω_z) tracking is *unexpectedly better in MuJoCo* (0.201 vs. 0.305). This partial reversal challenges the simple assumption that “training simulator is always better.”

The ω_z reversal is likely caused by MuJoCo’s joint damping (**damping=0.1**) acting as additional passive resistance that suppresses yaw oscillations. In Isaac Lab, joint damping is 0.0 (friction is handled by the actuator model), allowing more rotational jitter. This beneficial side-effect of MuJoCo’s extra damping improves angular tracking despite not being present during training.

Oscillation Amplification—Strongly supported. The velocity time series (Figure 3) reveal dramatically different oscillation characteristics:

- **Isaac Lab:** v_x oscillates tightly around the command with small amplitude (~ 0.2 m/s peak-to-peak). Reaches steady state within ~ 1 second.
- **MuJoCo:** v_x shows significantly larger oscillations (~ 0.8 m/s peak-to-peak), swinging between 0.2 and 1.0 m/s throughout the entire episode.

This oscillation is the *dominant source* of the RMSE gap and is consistent with the explicit PD controller in MuJoCo interacting with the contact model: the policy commands a position target $\rightarrow$ the explicit PD computes torque based on the *previous* timestep’s error $\rightarrow$ the one-step delay causes overshoot $\rightarrow$ contact forces create a correction $\rightarrow$ the cycle repeats at the gait frequency.

Contact Model Divergence in ω_z —Refuted. The largest per-scenario gap appears in the **Diagonal** scenario (v_x/v_y combined, $3.17\times$ ratio), not in ω_z . Angular velocity actually transfers *better* than linear velocities. The true bottleneck is actuator dynamics (implicit vs. explicit PD), not the contact/friction model.

Root cause: Actuator dynamics mismatch. The sim-to-sim gap is dominated by the **implicit vs. explicit PD actuator model**:

1. **Isaac Lab (implicit PD):** Position targets are integrated into the physics solver. The solver simultaneously resolves PD forces and contact forces, producing smooth, stable torque application.
2. **MuJoCo (explicit PD):** $\tau = K_p(q_{\text{target}} - q_{t-1}) - K_d\dot{q}_{t-1}$ is computed *before* the physics step using the *previous* state. This one-step delay creates a phase lag between the error signal and the corrective force, inherently prone to oscillation at the control rate.

This finding has important implications: **friction/mass randomization** (what the three DR levels primarily vary) addresses contact force differences but does *not* address the actuator dynamics gap. An **actuator network** approach—training a neural network to mimic MuJoCo’s actuator response—may be more effective at closing this specific gap than wider DR ranges.

2.10.6 Conclusion (Experiment 1B)

1. **Transfer Fidelity:** Supported for linear velocity ($2\text{--}3\times$ worse in MuJoCo). Reversed for angular velocity (ω_z is $0.66\times$ better in MuJoCo due to passive joint damping).
2. **Oscillation:** Strongly supported. MuJoCo shows $\sim 4\times$ larger velocity oscillation amplitude, driven by explicit PD one-step delay.
3. **Contact divergence:** Refuted. The dominant gap source is actuator dynamics, not contact/friction modeling.
4. **Sim-gap quantification:** 39% higher overall RMSE in MuJoCo (0.273 vs. 0.197), but 100% survival—the gap degrades precision without breaking stability.
5. **Recommendation:** Combine Moderate DR (for contact/mass robustness) with an actuator network (for dynamics fidelity).

2.11 Experiment 1C: ActuatorNet — Learned Actuator Model for Mismatch Reduction

2.11.1 Introduction and Motivation

Experiment 1B identified **actuator dynamics mismatch** (implicit vs. explicit PD integration) as the dominant source of the sim-to-sim gap, responsible for $\sim 4\times$ larger velocity oscillation in MuJoCo compared to Isaac Lab. Domain randomization (Experiment 1A) addresses contact and mass mismatch but does *not* target this actuator gap—DR randomizes friction, mass, and PD gains, but the policy still expects the implicit PD actuator response it was trained with.

The **ActuatorNet** approach [2] offers a targeted solution: instead of using an analytical PD controller ($\tau = K_p(q_{\text{des}} - q) - K_d\dot{q}$) in MuJoCo deployment, I train a *neural network* to learn the actual actuator dynamics from MuJoCo simulation data. This learned model captures nonlinear effects—contact-dependent friction, velocity-dependent damping, torque saturation, and joint-specific characteristics—that a simple PD formula cannot represent.

The key insight is that the policy was trained with Isaac Lab’s implicit PD actuator, which produces a specific torque response to position commands. If I can learn a *more accurate* mapping from position commands to torques in MuJoCo, the deployed policy “sees” actuator behavior closer to what it expects, reducing the transfer gap.

2.11.2 Methodology

The ActuatorNet pipeline consists of three stages: data collection, network training, and deployment integration.

Stage 1: Data Collection from MuJoCo. I collect actuator dynamics data by running the Go2 MuJoCo model through diverse joint configurations:

- **Phase A — Random Motion:** 200 episodes $\times$ 500 steps. Random joint position offsets (± 0.8 rad from default) are applied every 50 steps, producing natural dynamic trajectories that explore the typical operating range.
- **Phase B — Systematic Sweeps:** 20 sweeps per joint (12 joints). Each joint is swept from -1.5 to $+1.5$ rad in 100 steps, providing systematic coverage of the full joint range.

At each simulation step, I record the tuple $(j, q_{\text{des}}, q_{\text{current}}, \dot{q}, \tau_{\text{actual}})$ where τ_{actual} is extracted from MuJoCo’s `qfrc_actuator`, which includes the effects of `ctrlrange` clipping, joint damping, armature, and friction.

Table 19: ActuatorNet Training Data Statistics

Property	Value
Total samples	1,488,000 (1.49M)
Samples per joint	124,000 (perfectly balanced)
Position range (q_{des})	$[-2.30, +1.80]$ rad
Velocity range ($\dot{q}$)	$[-31.2, +35.8]$ rad/s
Torque range (τ)	$[-29.4, +45.4]$ N·m
File size	128 MB (CSV)

Stage 2: Network Architecture and Training. The ActuatorNet is a small MLP that predicts the torque for a single joint given the desired position, current state, and joint identity:

Table 20: ActuatorNet Architecture

Component	Specification
Input	15-dim: $[q_{\text{des}}, q_{\text{cur}}, \dot{q}, \text{one-hot}(j)]$
Hidden layers	[128, 64, 32] with LayerNorm + ELU + Dropout(0.1)
Output	1-dim: predicted torque $\hat{\tau}$
Loss function	MSE: $\mathcal{L} = \frac{1}{N} \sum_i (\hat{\tau}_i - \tau_i)^2$
Optimizer	Adam (lr= 10^{-3} , weight decay= 10^{-5})
Scheduler	ReduceLROnPlateau (patience=10, factor=0.5)
Train/val split	80% / 20% (random)
Epochs	100 (early stopping on validation loss)
Model size	71 KB (TorchScript)

The 12-dimensional one-hot encoding of the joint index allows the network to learn *joint-specific* dynamics: hip joints, thigh joints, and calf joints have different friction characteristics, torque limits, and inertial loading in MuJoCo. A shared network with joint conditioning is more parameter-efficient than training 12 separate models.

All input features are normalized using a **StandardScaler** fitted on the training set, which is critical for convergence given the wide range of velocities (± 35 rad/s) and torques (± 45 N·m).

Stage 3: Deployment Integration. During MuJoCo deployment, the ActuatorNet replaces the analytical PD controller in the control loop:

```
# Standard PD (Experiments 1A/1B):
tau = Kp * (q_target - q_current) - Kd * q_vel

# ActuatorNet replacement:
features = [q_target, q_current, q_vel, one_hot(joint_id)]
features_scaled = scaler.transform(features)
tau = actuator_net(features_scaled) # Learned torque prediction
```

Listing 2: ActuatorNet deployment: replaces PD with learned model

The rest of the deployment pipeline (observation construction, joint reordering, control swap) remains identical to Experiments 1A/1B, isolating the effect of the actuator model change.

2.11.3 Hypotheses

Hypothesis 7 (Oscillation Reduction). *The ActuatorNet will reduce the high-frequency velocity oscillation observed in MuJoCo (Experiment 1B), because the learned model better captures the actual torque response and avoids the one-step delay artifacts of the analytical PD controller.*

Hypothesis 8 (Multi-Axis Improvement). *The ActuatorNet will particularly improve multi-axis scenarios (Diagonal, Circle), where the analytical PD's oscillation interacts with multi-directional ground contact forces to create compound tracking errors.*

Hypothesis 9 (Learned Dynamics Trade-off). *The ActuatorNet may introduce new trade-offs: the learned model may be less accurate in operating regions poorly covered by the training data (e.g., extreme joint velocities during fast locomotion), potentially degrading some scenarios while improving others.*

2.11.4 Experiment Setup

- **Policy:** Baseline (same policy used in Experiments 1A/1B, trained with standard DCMotor actuator in Isaac Lab, 5,000 iterations)
- **Actuator model:** ActuatorNet (TorchScript, 71 KB) replaces PD controller
- **Evaluation:** Same 8 directional scenarios, 5 episodes $\times$ 500 steps each
- **Comparison:** Direct comparison with the same Baseline policy using analytical PD (from Experiment 1A)
- **Control variable:** Only the actuator model differs; observation construction, joint mapping, control swap, and policy are identical

2.11.5 Results

Survival and Success Rate. The ActuatorNet achieves **100% success rate** (40/40 episodes) across all 8 scenarios, matching the analytical PD baseline. The learned actuator model does not introduce instability.

Table 21: Per-Scenario RMSE: Analytical PD vs. ActuatorNet

Scenario	PD Baseline	ActuatorNet	Change
Forward	0.282	0.377	+33.7%
Backward	0.269	0.239	−11.2%
Left	0.249	0.216	−13.3%
Right	0.231	0.244	+5.6%
Turn_Left	0.287	0.288	+0.3%
Turn_Right	0.288	0.279	−3.1%
Diagonal	0.507	0.428	−15.6%
Circle	0.407	0.374	−8.1%
Average	0.315	0.306	−2.9%

Velocity Tracking RMSE.

Table 22: Per-Component Mean RMSE: PD vs. ActuatorNet

Component	PD Baseline	ActuatorNet	Change
RMSE v_x (mean)	0.231	0.238	+3.0%
RMSE v_y (mean)	0.209	0.186	−11.0%
RMSE ω_z (mean)	0.300	0.296	−1.3%

Per-Component RMSE Analysis. The ActuatorNet’s most significant improvement is in **lateral velocity (v_y) tracking**, with an 11% reduction in mean RMSE. The yaw rate (ω_z) also improves slightly (−1.3%), while forward velocity (v_x) is marginally worse (+3.0%).

Diagonal Scenario Deep Dive. The Diagonal scenario—identified as the hardest in Experiment 1A—shows the most dramatic improvement:

Table 23: Diagonal Scenario: PD vs. ActuatorNet (Component Breakdown)


Component	PD Baseline	ActuatorNet	Change
RMSE v_x	0.480	0.466	−2.9%
RMSE v_y	0.517	0.478	−7.5%
RMSE ω_z	0.522	0.323	− 38.1%
Overall	0.507	0.428	−15.6%

The ω_z component improves by **38.1%**—from 0.522 (PD) to 0.323 (ActuatorNet). This is precisely the yaw coupling problem that plagued the Aggressive DR policy in Experiment 1A (where Diagonal ω_z RMSE was 0.645). The ActuatorNet addresses this coupling more effectively than either DR strategy, suggesting that the yaw disturbance during diagonal locomotion is driven by actuator dynamics rather than contact forces.


Table 24: Locomotion Quality: PD vs. ActuatorNet

Metric	PD Baseline	ActuatorNet	Observation
Action smoothness	0.858	1.033	20% less smooth
Mean torque (N·m)	16.9	17.2	Similar (+1.8%)
Mean height (m)	0.323	0.306	1.7 cm lower


Locomotion Quality. The ActuatorNet produces **less smooth actions** (20% higher action rate) and achieves a **slightly lower stance height** (0.306 m vs. 0.323 m). The height reduction is scenario-dependent: lateral and turning scenarios maintain normal height ($\sim 0.35\text{--}0.37$ m), while forward-dominant scenarios (Forward, Diagonal, Circle) show significantly lower height (~ 0.22 m), suggesting the learned actuator model alters the forward gait strategy.

Go2 DR Comparison — Overall

(a) Overall comparison



(b) Per-scenario and per-component RMSE



(c) Velocity tracking timeseries

Figure 4: ActuatorNet vs analytical PD comparison. **(a)** Overall: ActuatorNet achieves 100% success with improved multi-axis tracking but height trade-off. **(b)** RMSE: excels on Diagonal/Circle (multi-axis), regresses on Forward. **(c)** Timeseries: reduces yaw oscillation but exhibits lower stance height.

2.11.6 Analysis and Discussion

Oscillation Reduction—Partially supported. The ActuatorNet reduces overall RMSE by 2.9% ($0.315 \rightarrow 0.306$), and improves lateral velocity tracking by 11%. However, the improvement is not uniform: forward velocity (v_x) tracking actually worsens slightly (+3%). The oscillation reduction is most effective in the *lateral* and *yaw* components, suggesting the learned model better captures the friction and contact dynamics relevant to sideways motion, but may slightly over-compensate in the sagittal (forward) plane.

Multi-Axis Improvement—Strongly supported. The multi-axis scenarios show the clearest benefit:

- **Diagonal:** RMSE -15.6% ($0.507 \rightarrow 0.428$), with ω_z component -38.1%
- **Circle:** RMSE -8.1% ($0.407 \rightarrow 0.374$)

The Diagonal ω_z improvement (38%) is particularly significant because this was the exact failure mode that degraded Aggressive DR in Experiment 1A. The ActuatorNet addresses the root cause (actuator-contact coupling during multi-axis motion) more effectively than wider DR ranges.

Learned Dynamics Trade-off—Supported. The ActuatorNet introduces a clear trade-off:

- **Improved:** Backward (−11.2%), Left (−13.3%), Diagonal (−15.6%), Circle (−8.1%)
- **Degraded:** Forward (+33.7%), Right (+5.6%)
- **Neutral:** Turn_Left (+0.3%), Turn_Right (−3.1%)

The Forward scenario degradation (+33.7%) is the largest regression and is accompanied by a significant height drop (0.351 m $\rightarrow$ 0.217 m). This suggests the ActuatorNet’s learned dynamics produce a different gait for forward locomotion—a lower, crouching gait that may result from the network slightly under-predicting the torques needed to maintain height during dynamic forward walking. This is consistent with the data collection strategy: random motions and systematic sweeps may under-represent the sustained high-torque profiles needed for forward trotting.

Comparison with Domain Randomization. The ActuatorNet and Moderate DR address different components of the sim-gap:

Table 25: ActuatorNet vs. Moderate DR: Complementary Approaches

Gap Source	Moderate DR	ActuatorNet
Contact/friction mismatch	Addresses directly	Does not address
Mass distribution	Addresses directly	Does not address
Actuator dynamics	Partially (gain randomization)	Addresses directly
Multi-axis coupling	Limited	Effective (−38% on Diagonal ω_z)

This suggests that combining both approaches—Moderate DR during training (for contact/mass robustness) with ActuatorNet during deployment (for actuator fidelity)—may yield the best overall transfer quality, a direction for future work.

2.11.7 Conclusion (Experiment 1C)

1. **Oscillation reduction:** Partially supported. Overall RMSE improved by 2.9%, with 11% improvement in lateral velocity. Forward velocity slightly degraded.
2. **Multi-axis improvement:** Strongly supported. Diagonal RMSE improved by 15.6%, with yaw component improving by 38.1%—directly addressing the dominant failure mode identified in Experiments 1A and 1B.
3. **Learned trade-off:** Supported. The ActuatorNet improves 5/8 scenarios but degrades Forward tracking by 33.7%, revealing insufficient training data coverage for forward locomotion. The learned model exhibits a gait strategy shift (lower stance) that harms straight-line walking.
4. **Practical finding:** The ActuatorNet is most effective where the analytical PD is weakest—multi-axis scenarios with complex actuator-contact coupling. It is a complementary approach to DR, targeting different mismatch sources.
5. **Future improvement:** Enriching the training data with sustained locomotion trajectories (not just random motion and sweeps) may improve the forward scenario and eliminate the height deficit.

2.12 Experiment 1D: Rapid Motor Adaptation (RMA) for Online Dynamics Estimation

2.12.1 Introduction and Motivation

Experiments 1A–1C addressed the sim-to-sim gap through training-time robustness (Domain Randomization) and deployment-time actuator modeling (ActuatorNet). Both approaches are

offline: DR creates a fixed robust policy, and ActuatorNet creates a fixed learned controller. Neither can *adapt* to the actual dynamics encountered during deployment.

Rapid Motor Adaptation (RMA) [4] offers a fundamentally different paradigm: the policy *online estimates* the current environment dynamics from observation history and conditions its behavior on this estimate. During training, the policy receives *privileged information* (ground-truth physics parameters) through an environment encoder. During deployment, an adaptation module replaces the encoder by predicting the latent dynamics representation from a buffer of recent observations—no privileged information required.

This approach is particularly promising for sim-to-sim transfer because the MuJoCo dynamics are *unknown but fixed*: the adaptation module can converge to a stable latent estimate within seconds, effectively “identifying” MuJoCo’s physics and adjusting the policy accordingly.

2.12.2 Architecture

The RMA framework consists of three modules trained in two phases:




Figure 5: RMA pipeline: privileged encoder at training is replaced by observation-history-based adaptation at deployment.

Table 26: RMA Module Architecture (Go2)

Module	Input	Architecture	Output
Env. Encoder	privileged info (4)	[32, 16], ELU	latent z (4)
Base Policy	obs (48) + z (4) = 52	[256, 256, 256], ELU	action (12)
Adaptation	obs history (50×48=2400)	[256, 128], LayerNorm, Dropout(0.1)	$\hat{z}$ (4)

2.12.3 Privileged Information

During Phase 1 training, the encoder receives 4-dimensional ground-truth physics parameters that the policy cannot observe through its standard sensors:

Table 27: Go2 RMA Privileged Information

Parameter	Dim	Range	Description
friction_scale	1	[0.5, 1.5]	Ground friction scale
mass_offset	1	[−0.5, +1.0] kg	Base mass offset (additive)
kp_scale	1	[0.8, 1.2]×	PD stiffness / nominal (20.0)
kd_scale	1	[0.8, 1.2]×	PD damping / nominal (0.5)

These ranges match the **Moderate DR** configuration from Experiment 1A, ensuring the encoder learns a meaningful latent representation over the same parameter space. The privileged information is normalized to $[-1, 1]$ before encoding.

2.12.4 Two-Phase Training

Phase 1: Base Policy with Privileged Encoder. The policy and encoder are trained jointly with PPO. The environment provides both the standard 48-dimensional observation and the 4-dimensional privileged information at each step. The encoder compresses the privileged information into a 4-dimensional latent z , which is concatenated with the observation to form the policy input (52 dimensions total).

Training configuration:

- Task: `Isaac-Velocity-Flat-Go2-Sim2Sim-RMA-v0`
- Network: [256, 256, 256] with ELU activation
- Iterations: 5,000
- Learning rate: 1.0×10^{-3}
- 4,096 parallel environments with Moderate DR

The encoder learns to extract task-relevant dynamics features: for instance, it may learn that low friction correlates with different contact dynamics, or that high mass requires larger torques for the same acceleration. The key property is that the encoder’s 4-dimensional output captures the *dynamics-relevant* projection of the physics parameters, discarding irrelevant variation.

Phase 2: Adaptation Module Training (Supervised). With the encoder and policy *frozen*, the adaptation module is trained to predict the encoder’s output from observation history alone:

$$\mathcal{L} = \text{MSE}(\hat{z}_{\text{adapt}}, z_{\text{encoder}}) = \frac{1}{N} \sum_i \|f_{\text{adapt}}(h_i) - f_{\text{encoder}}(e_i)\|^2 \quad (3)$$

where h_i is the flattened observation history ($50 \times 48 = 2400$ dimensions) and e_i is the privileged information.

Data collection: Two strategies are implemented:

1. **Isaac Lab collection** (`collect_rma_data.py`): Rollout the trained policy in Isaac Lab, recording (obs_history, privileged_info) pairs. 100,000 samples from 1,024 parallel environments.
2. **MuJoCo collection** (`collect_mujoco_rma_data.py`): Run the policy in MuJoCo with randomized physics, building observations in Isaac Lab order. This produces data that matches the *deployment distribution*, potentially improving adaptation quality.

The MuJoCo collection strategy is preferred for sim-to-sim transfer because the adaptation module trains on observations from the same simulator it will deploy in.

Training configuration:

- Epochs: 100 (with early stopping on validation loss)
- Batch size: 256
- Optimizer: AdamW (lr= 10^{-3} , weight decay= 10^{-5})
- Scheduler: ReduceLROnPlateau (patience=10, factor=0.5)
- Train/val split: 80%/20%

2.12.5 Deployment in MuJoCo

The deployment script (`deploy_rma_mujoco.py`) runs the complete RMA pipeline:

1. **Build observation** (48-dim) from MuJoCo state (identical to Experiment 1A pipeline)
2. **Update history buffer:** Rolling buffer of 50 most recent observations (1 second at 50 Hz)

3. **Adaptation step:** $\hat{z} = f_{\text{adapt}}(\text{flatten}(\text{history}))$ — predicts 4-dim latent from 2400-dim flattened history
4. **Policy step:** $a = \pi(\text{concat}(\text{obs}, \hat{z}))$ — 52-dim input, 12-dim output
5. **Apply control:** Convert actions to PD targets and step physics (identical to Experiment 1A)

Online adaptation: Unlike the fixed policies of Experiments 1A–1C, the RMA policy’s behavior *evolves* during the episode. In the first ~ 1 second (50 steps), the history buffer fills and the latent estimate stabilizes. After convergence, the policy operates with a dynamics estimate that reflects MuJoCo’s actual physics, potentially producing more appropriate torque profiles than a fixed policy trained with average DR parameters.

2.12.6 Comparison with Other Approaches

Table 28: Transfer Strategy Comparison

Property	Standard	DR	ActuatorNet	RMA
Trains with privileged info	No	No	No	Yes
Online adaptation	No	No	No	Yes
Addresses actuator gap	No	Partially	Yes	Potentially
Addresses contact gap	No	Yes	No	Potentially
Requires MuJoCo data	No	No	Yes	Optional
Additional deployment cost	None	None	71 KB model	2400-dim history

2.12.7 Phase 1 Evaluation Results

The Phase 1 policy (with privileged observations) was trained successfully for 5,000 iterations. To understand the policy’s behavior before completing Phase 2 adaptation module training, I evaluated it using *dummy privileged observations* (friction=1.0, mass_offset=0.0, kp_scale=1.0, kd_scale=1.0)—representing an unrandomized environment.

Evaluation Protocol. Two test conditions were examined:

1. **Flat terrain** (training environment): MuJoCo `scene_isaacclab_terrain.xml` with flat ground
2. **Rough terrain** (out-of-distribution): MuJoCo `scene_terrain.xml` with procedurally-generated heightfield

Each scenario (Forward, Backward, Left, Right, Turn_Left, Turn_Right, Diagonal, Circle) was run for 5 episodes of 500 steps (10 seconds at 50 Hz). Success criterion: robot survives $\geq 80\%$ of episode duration without falling (base height < 0.15 m).

Table 29: RMA Phase 1 Performance with Dummy Privileged Observations

Scenario	Flat Terrain			Rough Terrain		
	Success	RMSE vx	Steps	Success	RMSE vx	Steps
Forward	100%	0.482	500	100%	0.483	500
Backward	0%	0.371	286	100%	0.406	500
Left	100%	0.051	500	100%	0.036	500
Right	0%	0.320	52	100%	0.032	500
Turn_Left	100%	0.008	500	100%	0.005	500
Turn_Right	100%	0.006	500	100%	0.005	500
Diagonal	0%	0.240	67	0%	0.399	255
Circle	100%	0.471	500	100%	0.461	500
Overall	62.5%	0.244	—	87.5%	0.228	—

Results.

Analysis: Terrain Effect. The RMA Phase 1 policy performs *better* on rough terrain (87.5% success) than on flat terrain (62.5% success). Three scenarios (Backward, Right) that fail immediately on flat terrain succeed on rough terrain. This behavior is expected given the testing setup:

1. **Dummy privileged values provide no information:** Fixed values (friction=1.0, mass=0.0, etc.) are constants, not estimates. The policy was trained to condition on privileged info, but feeding it uninformative constants produces undefined behavior.
2. **Terrain geometry provides stabilization:** The rough heightfield may prevent certain fall modes (e.g., sideways tipping in Right scenario) by providing lateral contact support that flat ground lacks.
3. **Training distribution mismatch:** The policy was trained with moderate DR (friction 0.5–1.5, mass ± 0.5 kg), but deployed on perfect flat ground with dummy values representing “nominal” physics. The rough terrain’s contact variability may paradoxically be *closer* to the training distribution than the static flat environment.

Comparison with Moderate DR Baseline. For context, the Moderate DR policy (Experiment 1A) achieves:

- **Flat terrain:** 100% success, Forward RMSE = 0.108 m/s
- **Rough terrain:** 100% success, but 94% distance reduction

The RMA Phase 1 policy’s Forward RMSE (0.482 m/s) is **4.5 $\times$ worse** than Moderate DR. This confirms that Phase 1 alone is *not deployable*—it requires the Phase 2 adaptation module to function properly.

2.12.8 Critical Finding: Phase 1 Cannot Deploy Without Phase 2

The evaluation reveals a fundamental limitation: **RMA Phase 1 policies trained with privileged observations cannot be deployed directly.** The policy learned to *rely on* the 4-dimensional latent z during training, but without the adaptation module to estimate z from observation history, the policy receives meaningless dummy values and produces suboptimal behavior.

This is not a flaw—it is by design. The RMA architecture explicitly assumes Phase 2 will be completed before deployment. However, for research environments where Phase 2 training is

resource-constrained, this dependency creates a deployment gap.

2.12.9 Guidelines for Future RMA Phase 2 Development

To complete the RMA pipeline and achieve deployable performance, the following steps are recommended:

1. Adaptation Module Training Data Collection. The adaptation module requires (observation_history, privileged_info) pairs collected from the trained Phase 1 policy. Two strategies:

- **Isaac Lab collection** (faster, cleaner): Rollout the policy in Isaac Lab’s training environment, recording 50-step observation windows and corresponding privileged parameters. Target: 100,000 samples from 1,024 parallel environments (~30 minutes).
- **MuJoCo collection** (preferred for sim-to-sim): Run the policy in MuJoCo with randomized physics parameters (friction, mass, PD gains varied within Moderate DR ranges), collecting observation histories. This produces data matching the *deployment distribution*, potentially improving adaptation quality. Target: 50,000 samples (~2 hours with parallel workers).

Recommendation: Use MuJoCo collection to avoid distribution shift. The adaptation module will train on observations from the same simulator it deploys in, capturing MuJoCo-specific artifacts (e.g., contact dynamics, numerical integration differences).

2. Adaptation Module Architecture Tuning. The current architecture (2400-dim input $\rightarrow$ [256, 128] $\rightarrow$ 4-dim output) may be oversized for the task. Consider:

- **Reduce history length:** 50 steps (1 second) may be excessive. Test 25 steps (0.5s) or 10 steps (0.2s) to reduce input dimensionality and computational cost.
- **Use temporal convolutions:** Replace flattened history with 1D CNN over time axis to capture temporal patterns (e.g., oscillation frequency, settling time) that encode dynamics.
- **Attention mechanism:** Use self-attention over the observation sequence to let the model focus on task-relevant time windows (e.g., contact events, rapid accelerations).

3. Loss Function Enhancement. The standard MSE loss (Equation 3) treats all latent dimensions equally. Consider:

- **Weighted MSE:** If certain privileged parameters (e.g., friction) are more important for policy performance than others (e.g., kd_scale), weight the loss accordingly: $\mathcal{L} = \sum_i w_i (\hat{z}_i - z_i)^2$.
- **Task-aware loss:** Instead of matching encoder output, directly optimize for task performance. Augment loss with rollout-based term: $\mathcal{L}_{\text{total}} = \mathcal{L}_{\text{MSE}} + \lambda \mathcal{L}_{\text{task}}$, where $\mathcal{L}_{\text{task}}$ penalizes poor tracking accuracy or instability when using predicted $\hat{z}$.

4. Evaluation Protocol for Phase 2. After training the adaptation module, evaluate on the same 8 scenarios (flat + rough terrain) and compare against:

1. **Moderate DR (baseline):** 100% flat success, 0.108 m/s RMSE
2. **Phase 1 with dummy obs:** 62.5% flat success, 0.482 m/s RMSE
3. **Phase 2 with adaptation:** Target $\geq 95\%$ success, ≤ 0.15 m/s RMSE

Key metrics:

- **Success rate improvement:** Phase 2 should recover Backward, Right scenarios that failed in Phase 1.
- **Tracking accuracy:** Forward RMSE should approach Moderate DR (0.108 m/s).
- **Adaptation convergence time:** Measure how quickly $\hat{z}$ stabilizes (ideally $<0.5s$).
- **Rough terrain generalization:** Does adaptation help on out-of-distribution terrain?

5. Alternative: Simplified RMA for Sim-to-Sim. If Phase 2 training proves too costly, consider a lightweight alternative:

- **Fixed latent sweep:** Instead of training an adaptation module, sweep through a grid of fixed z values (e.g., $z \in \{-1, 0, +1\}^4$) during deployment and select the one that maximizes a reward proxy (e.g., velocity tracking error over 1 second). This requires no additional training but loses the online adaptation benefit.
- **System identification at startup:** Run a 5-second “calibration phase” where the robot executes random actions and observes responses, then uses least-squares to fit MuJoCo’s PD parameters. Pass these estimates as privileged info. This is a classical control approach that bypasses the need for learned adaptation.

6. When to Skip RMA. RMA is most valuable when:

- Deployment environment dynamics are *unknown and variable* (e.g., different robots, changing payloads, worn actuators)
- Online adaptation provides measurable benefit over fixed robust policies
- Training resources (time, compute, data) for Phase 2 are available

For sim-to-sim transfer where MuJoCo physics are *fixed and known*, RMA may be overkill:

- **Moderate DR already achieves 100% flat success** with better tracking (0.108 vs 0.482 m/s)
- **No dynamics variation during deployment**—MuJoCo is deterministic
- **ActuatorNet addresses the actuator gap more directly** (Experiment 1C)

Recommendation for this experiment: Given time constraints and strong Moderate DR baseline results, **deprioritize RMA Phase 2 completion**. The Phase 1 evaluation already provides valuable insights (privileged obs dependency, terrain effects). Moderate DR alone achieves 100% success on flat terrain (0.108 RMSE) and is recommended over ActuatorNet for general deployment given ActuatorNet’s Forward scenario degradation (Experiment 1C).

2.13 Experiment 1E: Terrain Robustness — Flat vs. Rough Terrain Transfer

2.13.1 Introduction and Motivation

Having established baseline transfer performance on flat terrain (Experiments 1A–1D), now evaluate whether the trained policies can generalize to **rough terrain**—a geometric complexity not present during training. This tests whether parameter-based Domain Randomization (friction, mass, PD gains) provides robustness to terrain geometry, or whether explicit terrain variation during training is required for effective navigation. This tests the generalization capability of the policies to out-of-distribution terrain geometries, a key challenge in robotic locomotion [8].

Experiments 1A–1D established that while flat-terrain policies can transfer to MuJoCo, their performance is sensitive to the level of Domain Randomization (DR). In this experiment, I escalate the challenge by evaluating policies on **rough terrain** in MuJoCo. I compare two classes of policies trained in Isaac Lab:

1. **Flat-trained policies:** Baseline, Moderate, and Aggressive DR policies trained exclusively on flat ground.
2. **Rough-trained policies:** Baseline and Moderate DR policies trained on Isaac Lab’s procedurally generated rough terrain (heightfields with vertical variations of ± 0.1 m).

This comparison identifies whether geometric robustness can emerge from parameter DR alone, or if explicit terrain variation during training is required for effective navigation.

2.13.2 Hypotheses

Hypothesis 10 (Terrain Generalization via DR). *For flat-trained policies, higher DR levels (Moderate) will marginally improve distance traveled on rough terrain compared to Baseline, as randomization of contact parameters provides a narrow proxy for geometric variation.*

Hypothesis 11 (Training Morphology Mismatch). *Rough-trained policies will achieve significantly higher forward progress (meters traveled) on rough terrain than flat-trained policies, as they have learned gait patterns specifically adapted to geometric irregularities.*

Hypothesis 12 (Survival-Progress Trade-off). *While rough-trained policies may move faster, they will exhibit lower survival rates in MuJoCo compared to Isaac Lab, as the sim-to-sim physics gap (actuator delay, contact solver) is amplified by the aggressive gait required for rough terrain.*

2.13.3 Experiment Setup

- **Policies:**
 - **Flat-trained:** Baseline, Moderate, Aggressive (Experiment 1A)
 - **Rough-trained:** Rough_Baseline, Rough_Moderate
- **Terrain:** scene_isaaclab_terrain.xml (8 m $\times$ 2 m heightfield; vertical range: ± 0.1 m)
- **Command:** Forward velocity $v_x = 0.5$ m/s
- **Evaluation:** 3 trials per policy/terrain, 20 seconds max duration
- **Metrics:**
 - Success rate (survival without falling)
 - Distance traveled (maximum Euclidean displacement from start)
 - Velocity tracking RMSE (v_x, v_y, ω_z)
 - Height variance (base height stability)
 - Tilt metric (XY component of gravity vector)

2.13.4 Results

Table 30: Terrain Comparison: Success Rate and Distance Traveled on Rough Terrain

Training	Policy	Rough Success	Rough Dist (m)	Avg Survival (s)
Flat	Baseline	100% (3/3)	0.94	20.0
Flat	Moderate	100% (3/3)	1.09	20.0
Flat	Aggressive	100% (3/3)	0.51	20.0
Rough	Baseline	100% (3/3)	3.41	20.0
Rough	Moderate	100% (3/3)	4.15	20.0

Success Rate and Survival. All five policies achieve **100% success rate** on rough terrain, demonstrating successful sim-to-sim transfer. However, the distance traveled reveals a dramatic distinction between training approaches:

Flat-trained policies show severe performance degradation on rough terrain:

- **Baseline:** 95% distance reduction (17.95 m $\rightarrow$ 0.94 m)
- **Moderate:** 94% distance reduction (17.78 m $\rightarrow$ 1.09 m)
- **Aggressive:** 97% distance reduction (18.02 m $\rightarrow$ 0.51 m)

Rough-trained policies demonstrate the critical importance of terrain-specific training:

- **Rough_Baseline:** Achieves **3.41 m** distance—a $3.6\times$ improvement over the best flat-trained policy (Moderate, 1.09 m)
- **Rough_Moderate:** Achieves **4.15 m** distance—a $3.8\times$ improvement over flat-trained Moderate, demonstrating that DR combined with terrain training provides the best rough terrain performance

Among flat-trained policies, Moderate DR achieves the highest distance (1.09 m, +16% over Baseline), while Aggressive degrades the most (0.51 m, -46% from Baseline). This confirms that moderate DR provides useful robustness, but excessive DR harms terrain adaptation.

Table 31: Velocity Tracking and Stability Metrics on Rough Terrain

Training	Policy	Rough v_x RMSE	Rough Z-Var	Avg Tilt
Flat	Baseline	0.476	0.00099	0.165
Flat	Moderate	0.471	0.00038	0.310
Flat	Aggressive	0.483	0.00171	0.326
Rough	Baseline	0.287	0.00089	0.080
Rough	Moderate	0.265	0.00062	0.075

Velocity Tracking Error. Forward velocity tracking shows a clear distinction between training approaches: flat-trained policies exhibit poor tracking (RMSE 0.471–0.483), while rough-trained policies achieve significantly better performance. **Rough_Moderate** achieves the best tracking (RMSE **0.265**), followed by Rough_Baseline (0.287)—both substantially better than any flat-trained policy. Height variance (Z-Var) is lowest for Rough_Moderate (0.00062) and Rough_Baseline (0.00089), demonstrating superior stability on uneven terrain compared to flat-trained policies.

Table 32: Body Tilt: Flat vs. Rough Terrain

Policy	Flat Tilt	Rough Tilt	Change
Baseline	0.029	0.165	$+5.7\times$
Moderate	0.031	0.310	$+10.0\times$
Aggressive	0.037	0.326	$+8.8\times$

Tilt and Stability. Body tilt increases $5.7\text{--}10\times$ on rough terrain. Surprisingly, **Baseline** shows the lowest tilt on rough terrain (0.165), while Moderate and Aggressive exhibit more aggressive body orientation (0.310–0.326). This suggests that DR-trained policies adopt more dynamic balance strategies that tolerate higher tilt angles, potentially trading conservative stance for forward progress.




Figure 6: Terrain robustness comparison: flat vs rough terrain across DR levels. **(a)** All policies achieve 100% success on both terrains. **(b)** Rough-trained policies achieve 3.6–3.8 $\times$ better distance than flat-trained (4.15 m vs 1.09 m); among flat-trained, Moderate DR best (+16% vs Baseline). **(c)** Rough-trained policies show superior velocity tracking (RMSE 0.265–0.287) compared to flat-trained (0.471–0.483). **(d)** Rough-trained policies (bottom) demonstrate effective obstacle navigation; flat-trained policies show conservative freezing behavior.

Hypothesis Evaluation. Terrain Generalization via DR: *Supported but limited.* For flat-trained policies, Moderate DR achieves the highest distance (1.09 m), but overall progress remains extremely low (<6% of commanded speed). Parameter DR alone is insufficient for rough terrain locomotion.

Training Morphology Mismatch: *Strongly supported.* Rough-trained policies dramatically outperform flat-trained policies: Rough_Baseline achieves **3.41 m** (3.6 $\times$ improvement over flat Moderate), while Rough_Moderate achieves **4.15 m** (3.8 $\times$ improvement) with the best tracking performance (RMSE 0.265). Explicit terrain training is the dominant factor for successful rough terrain transfer.

Survival-Progress Trade-off: *Not supported.* Contrary to the hypothesis, rough-trained policies achieve **100% success rate** while simultaneously moving significantly better than flat-trained policies. This demonstrates successful sim-to-sim transfer for rough terrain locomotion—both survival and performance are maintained. The combination of terrain-specific training and moderate DR (Rough_Moderate) provides the best overall performance: 4.15 m distance, 0.265 RMSE, and 100% survival.

Locomotion Strategy: Conservative vs. Adaptive. The results reveal two distinct locomotion strategies on rough terrain:

1. **Flat-trained policies (Conservative):** These policies prioritize stability over progress. When encountering unexpected geometric perturbations, they adopt a cautious stance that prevents falling but fails to clear obstacles effectively, resulting in minimal forward progress (~ 0.05 m/s achieved vs 0.5 m/s commanded). They achieve 100% survival by essentially "freezing" in place.

2. **Rough-trained policies (Adaptive):** These policies have learned gait patterns specifically adapted to geometric irregularities. They successfully "power through" obstacles while maintaining stability, achieving both superior translation (Rough_Moderate: 0.265 RMSE, 4.15 m distance) and 100% survival. The terrain-specific training enables robust rough terrain locomotion that transfers successfully to MuJoCo.

Implications for Sim-to-Real. This experiment demonstrates that **task-specific training is essential for robust transfer**. Policies trained on the target task (rough terrain) achieve successful sim-to-sim transfer with both high performance and stability. While flat-trained policies can survive on rough terrain, they fail to accomplish the locomotion task. The success of rough-trained policies validates that domain randomization combined with appropriate training environments can bridge the sim-to-sim gap effectively.

2.13.5 Conclusion (Experiment 1E)

I compared flat-trained and rough-trained Go2 policies on rough terrain in MuJoCo. Key findings:

1. **Training terrain is paramount:** Rough-trained policies achieve $3.6\text{--}3.8\times$ more distance than flat-trained policies (4.15 m vs 1.09 m for moderate DR), demonstrating that explicit terrain variation during training is essential for rough terrain performance.
2. **Successful rough terrain transfer:** Rough-trained policies achieve **100% success rate** in MuJoCo while simultaneously delivering superior locomotion performance (RMSE 0.265–0.287 vs 0.471–0.483 for flat-trained). This validates that task-specific training enables robust sim-to-sim transfer.
3. **DR amplifies terrain-specific training:** Rough_Moderate (terrain + DR) outperforms Rough_Baseline (terrain only) by 22% in distance (4.15 m vs 3.41 m) and 8% in tracking (RMSE 0.265 vs 0.287), confirming that moderate DR provides additional robustness even with terrain training.
4. **Locomotion strategies:** Flat-trained policies adopt conservative "freezing" behavior (100% survival, minimal progress), while rough-trained policies execute adaptive gaits that successfully navigate obstacles with both high performance and stability.

This experiment completes the Go2 evaluation suite, demonstrating that task-specific training combined with moderate domain randomization enables successful sim-to-sim transfer across both flat and geometrically complex terrains.

2.14 Go2 Transfer Summary: Key Findings and Lessons Learned

Having completed five experiments across flat and rough terrain, I synthesize the key findings for Go2 sim-to-sim transfer from Isaac Lab (PhysX) to MuJoCo.

2.14.1 Transfer Strategy Comparison

Table 33: Go2 Transfer Strategies: Trade-offs and Performance

Strategy	Approach	Benefits	Limitations
Baseline	No enhancements	Simple, fast	Narrow robustness
Moderate DR	Train-time randomization	Robust to contact variation	Doesn't address actuator gap
Aggressive DR	High randomization	Maximum contact robustness	Multi-axis coordination loss
ActuatorNet	Learned PD controller	Targets actuator gap	Requires MuJoCo data
RMA	Online adaptation	Adapts to deployment	Complex training

2.14.2 Key Insights

1. The "Goldilocks" Principle for Domain Randomization. Experiments 1A and 1E both demonstrate that **Moderate DR is optimal**:

- **Flat terrain:** Moderate = Baseline RMSE (0.257), Aggressive degrades to 0.341 (+33%)
- **Rough terrain:** Moderate achieves 1.09 m distance (+16% over Baseline), Aggressive degrades to 0.51 m (−46%)
- **Diagonal scenario:** Moderate ω_z RMSE = 0.371, Aggressive = 0.645 (+74%)

Excessive DR sacrifices multi-axis coordination for single-axis robustness, harming both sim-to-sim transfer and terrain generalization.

2. Dominant Mismatch Source: Actuator Dynamics. Experiment 1B identified the sim-to-sim gap:

- Isaac Lab (implicit PD): 1-step integration, constraint-based, very stiff
- MuJoCo (explicit PD): Torque = $K_p(q_d - q) + K_d(0 - \dot{q})$, applied as external force
- Result: 4× larger velocity oscillation in MuJoCo, 39% RMSE increase

DR randomizes friction, mass, and PD gains but does *not* address the implicit-vs-explicit PD gap. ActuatorNet directly targets this gap with a learned model.

3. Parameter DR $\neq$ Geometric Robustness. Experiment 1E revealed that parameter randomization provides limited terrain robustness:

- All policies: 94–97% distance reduction on rough terrain (17.8 m $\rightarrow$ 0.5–1.1 m)
- Height variance increases 13–43× on rough terrain
- Moderate DR provides +16% improvement over Baseline, but absolute performance remains poor

Lesson: Real-world deployment requires training on varied terrain *geometry* (stairs, slopes, gaps), not just randomized contact parameters.

4. Complementary Strategies. The three enhancement strategies address different aspects of the sim-to-sim gap:

- **DR**: Robustness to contact parameter variation (friction, mass) → helps with terrain roughness
- **ActuatorNet**: Bridges actuator dynamics mismatch → reduces oscillation, improves lateral tracking
- **RMA**: Online adaptation to deployment dynamics → complements both DR and ActuatorNet

Combining these strategies (e.g., Moderate DR + ActuatorNet + RMA) may provide the best overall transfer quality.

2.14.3 Lessons for Cassie and Flamingo

The Go2 experiments establish principles that guide the more complex Cassie (Part 2) and Flamingo (Part 3) transfers:

1. **Joint ordering and control swap are critical.** The Go2’s initial complete failure (Phase 1) was due to PD gain mismatch ($K_p = 25$ vs. 20) and missing control swap. Cassie and Flamingo have more complex joint mappings.
2. **Standing height validates physics match.** The Go2’s 0.06 m height difference (Isaac Lab: 0.40 m, MuJoCo: 0.34 m) reflects contact solver differences. Cassie shows larger mismatch due to passive ankles.
3. **Actuator dynamics dominate the sim-gap for position-controlled robots.** ActuatorNet’s success on Go2 suggests similar approaches may help Cassie. Flamingo’s hybrid control adds complexity.
4. **Moderate DR is a safe default.** Unless robot-specific analysis suggests otherwise, start with Moderate DR (friction: 0.8–1.2, mass: $\pm 15\%$, gains: $\pm 20\%$).

3 Part 2: Agility Cassie Biped

3.1 Robot Overview

Cassie [1] is a 3D bipedal robot developed by Agility Robotics, featuring a unique leg design inspired by ostrich anatomy. Unlike the Go2’s straightforward kinematic chain, Cassie’s legs incorporate a **4-bar linkage mechanism** that constrains the ankle joint, creating a fundamental structural mismatch between training and deployment simulators. Previous work on Cassie has explored various RL-based locomotion strategies for robust sim-to-sim and sim-to-real transfer [5, 13].

Table 34: Cassie Robot Specifications

Property	Value
Type	3D bipedal robot
Isaac Lab DOF	12 (6 per leg: hip abd., hip rot., hip flex., knee, ankle, toe)
MuJoCo actuators	10 (ankle joints are passive due to 4-bar linkage)
Leg mechanism	4-bar linkage (equality constraints in MuJoCo)
PD gains	Variable: hips $K_p = 100$, knees $K_p = 200$, toes $K_p = 20$
Mass	~ 32 kg
Standing height	~ 0.9 m



Figure 7: Agility Cassie biped robot. Note the complex leg mechanism with 4-bar linkage connecting the shin, tarsus, and ankle components.

Why is Cassie harder than Go2? The Go2 transfer required only joint reordering and gain matching—a one-to-one mapping between 12 trained joints and 12 deployed actuators. Cassie introduces three additional layers of complexity:

1. **Joint count mismatch:** 12 trained $\rightarrow$ 10 deployable (2 passive)
2. **Default pose mismatch:** Up to 0.6 rad (34.4°) difference in default joint angles
3. **Constraint physics:** The 4-bar linkage requires sub-millisecond timesteps for stable constraint solving

3.2 Transfer Challenges Unique to Cassie

3.2.1 Challenge 1: Joint Count Mismatch (12 $\rightarrow$ 10)

Isaac Lab models Cassie with 12 actuated joints—6 per leg. MuJoCo’s Cassie model has only 10 actuators because the `ankle_joint_left` and `ankle_joint_right` (Isaac Lab indices 4 and 10) are **passive**, constrained by the 4-bar linkage through equality constraints in the MuJoCo XML.

Table 35: Cassie Joint Mapping: Isaac Lab $\rightarrow$ MuJoCo

Isaac Idx	Joint Name	MuJoCo Actuator	Status
0	hip_abduction_left	left-hip-roll (act. 0)	Active
1	hip_rotation_left	left-hip-yaw (act. 1)	Active
2	hip_flexion_left	left-hip-pitch (act. 2)	Active
3	thigh_joint_left	left-knee (act. 3)	Active
4	ankle_joint_left	PASSIVE (4-bar)	No actuator
5	toe_joint_left	left-foot (act. 4)	Active
6	hip_abduction_right	right-hip-roll (act. 5)	Active
7	hip_rotation_right	right-hip-yaw (act. 6)	Active
8	hip_flexion_right	right-hip-pitch (act. 7)	Active
9	thigh_joint_right	right-knee (act. 8)	Active
10	ankle_joint_right	PASSIVE (4-bar)	No actuator
11	toe_joint_right	right-foot (act. 9)	Active

Impact: The policy outputs 12-dimensional actions, but actions at indices 4 and 10 cannot be applied—they are simply dropped during deployment. The policy wastes 2 action dimensions on joints it cannot control.

Observation workaround: For the ankle joint *observation* (needed by the policy as input), I use the `tarsus` joint position from MuJoCo as a proxy, since the tarsus is kinematically linked to the ankle through the 4-bar mechanism.

3.2.2 Challenge 2: Default Pose Mismatch

Isaac Lab and MuJoCo Cassie models have significantly different default (“home”) joint angles for the same physical standing pose:

Table 36: Cassie Default Pose Comparison

Joint	Isaac Lab	MuJoCo Home	Difference	Magnitude
hip_abduction_L	+0.10	0.00	−0.10	Small
hip_rotation_L	0.00	0.00	0.00	None
hip_flexion_L	+1.00	+0.50	− 0.50	28.6°
thigh_joint_L	−1.80	−1.20	+ 0.60	34.4°
ankle_joint_L	+1.57	<i>passive</i>	N/A	N/A
toe_joint_L	−1.57	−1.52	+0.05	Small

The most significant differences are in **hip_flexion** (0.5 rad) and **thigh_joint/knee** (0.6 rad). These correspond to an entirely different leg configuration: MuJoCo’s default is more upright with straighter legs, while Isaac Lab’s default has more bent knees.

Solution: Use MuJoCo’s home keyframe angles (not Isaac Lab defaults) as the observation baseline during deployment. Joint position observations are computed as $q_{\text{obs}} = q_{\text{current}} - q_{\text{mj_default}}$ rather than $q_{\text{current}} - q_{\text{isaac_default}}$.

3.2.3 Challenge 3: 4-Bar Linkage and Timestep Coupling

Cassie’s 4-bar linkage is implemented in MuJoCo via *equality constraints* with solver parameters `solref="0.005 1"`. The `solref` time constant (0.005 s) determines how many simulation steps are needed to resolve the constraint:

$$N_{\text{steps}} = \frac{t_{\text{solref}}}{\Delta t} = \frac{0.005}{0.0005} = 10 \text{ steps (at native dt)} \quad (4)$$

Critical finding: At the standard $\Delta t = 0.005$ s (200 Hz, matching Go2), the constraint solver gets only *1 step* to resolve—completely insufficient for the 4-bar linkage. This caused the constraints to violate immediately, the linkage to explode, and the robot to collapse.

Solution: Use MuJoCo’s native timestep of $\Delta t = 0.0005$ s (2 kHz, as specified in the Cassie XML) and adjust the control decimation to maintain the 50 Hz policy rate:

$$\text{control_decimation} = \frac{0.005 \times 4}{0.0005} = 40 \text{ substeps per policy step} \quad (5)$$

3.2.4 Challenge 4: Variable PD Gains and Motor Model

Unlike Go2’s uniform gains ($K_p = 20$, $K_d = 0.5$ for all joints), Cassie uses **joint-dependent** PD gains:

Table 37: Cassie PD Gains by Joint Type

Joint Type	K_p	K_d	Effort Limit (N·m)
hip_abduction	100.0	3.0	200
hip_rotation	100.0	3.0	200
hip_flexion	200.0	6.0	200
thigh_joint (knee)	200.0	6.0	200
ankle_joint	200.0	6.0	200 (passive in MuJoCo)
toe_joint	20.0	1.0	20

Furthermore, MuJoCo’s Cassie actuators use a **motor type** where `ctrl = torque / gear`. The gear ratios differ by joint (hip: gear=25, knee: gear=16, foot: gear=50), creating additional mismatch between the Isaac Lab effort limits (uniform 200 N·m) and MuJoCo’s effective torque range (e.g., hip-roll max = $25 \times 4.5 = 112.5$ N·m vs. Isaac Lab’s 200 N·m).

3.3 Training Configuration

3.3.1 Environment Configuration

Cassie training uses a **DirectRL environment** (not manager-based) for faster iteration, with the following specifications:

Table 38: Cassie Training Environment Configuration

Parameter	Value
Environment type	DirectRLEnv (custom Cassie)
Physics timestep	0.005 s (200 Hz)
Decimation	4 (policy at 50 Hz)
Episode length	20.0 s
Num environments	4,096
Observation dimension	48 (no scaling, raw values)
Action dimension	12 (position offsets, scale = 0.5)
Action scale	0.5 ($2 \times$ Go2’s 0.25)

Key difference from Go2: Cassie uses **no observation scaling**—raw sensor values are fed directly to the policy (no $\times 2.0$ for linear velocity, no $\times 0.25$ for angular velocity). This means the policy must internally learn to handle the different magnitudes of velocity vs. joint angle signals.

3.3.2 Reward Function

The Cassie reward function differs from Go2 in several important ways: higher penalties for torque usage, explicit joint deviation penalties for hip and toe joints, and a large termination penalty to discourage falls.

Table 39: Cassie Reward Function

Term	Weight	Purpose
track_lin_vel_xy_exp	+2.0	Exponential linear velocity tracking
track_ang_vel_z_exp	+1.0	Exponential angular velocity tracking
feet_air_time	+2.5	Encourage dynamic walking gait
lin_vel_z_l2	-2.0	Penalize vertical bouncing
ang_vel_xy_l2	-0.05	Penalize body roll/pitch rate
dof_torques_l2	-5.0×10^{-6}	Torque minimization
dof_acc_l2	-3.75×10^{-7}	Joint acceleration smoothness
action_rate_l2	-0.015	Action smoothness
flat_orientation_l2	-2.5	Keep upright orientation
joint_deviation_hip	-0.2	Penalize hip deviation from default
joint_deviation_toes	-0.2	Penalize toe deviation from default
dof_pos_limits	-1.0	Joint limit violation avoidance
feet_slide	-0.1	Penalize foot sliding in contact
termination_penalty	-200.0	Large penalty for falling

3.3.3 Observation Space (48 dimensions, unscaled)

Table 40: Cassie Observation Space (48 dimensions)

Index	Name	Dim	Scale
0-2	base_lin_vel (body frame)	3	$\times 1.0$ (raw)
3-5	base_ang_vel (body frame)	3	$\times 1.0$ (raw)
6-8	projected_gravity	3	$\times 1.0$
9-11	velocity_commands	3	$\times 1.0$
12-23	joint_positions (relative)	12	$\times 1.0$ (raw)
24-35	joint_velocities	12	$\times 1.0$ (raw)
36-47	previous_actions	12	$\times 1.0$

3.3.4 Network Architecture

Table 41: Cassie Policy Network Architecture

Component	Specification
Actor	$48 \rightarrow [256, 256, 256] \rightarrow 12$ (ELU activations)
Hidden layers	$[256, 256, 256]$ (vs. Go2's $[512, 256, 128]$)
Training iterations	30,000 ($6 \times$ Go2's 5,000)
Learning rate	$3.0e-4$ (vs. Go2's $1.0e-3$)
Empirical normalization	No

3.3.5 Domain Randomization

Three DR levels were trained for Cassie, following the same Baseline/Moderate/Aggressive structure as Go2 but with ranges tailored to Cassie's dynamics:

Table 42: Cassie Domain Randomization Configurations

Parameter	Baseline	Moderate	Aggressive
Friction	[0.5, 1.25]	[0.4, 1.35]	[0.3, 1.5]
Mass offset	$[-1.0, +2.0]$ kg	$[-1.5, +2.5]$ kg	$[-2.0, +3.0]$ kg
Gain scaling	$[0.8, 1.2] \times$	$[0.7, 1.3] \times$	$[0.6, 1.4] \times$
Joint position	± 0.5 rad	± 0.6 rad	± 0.75 rad
Push velocity	± 0.5 m/s	± 0.6 m/s	± 0.75 m/s
Push interval	10–15 s	8–13 s	6–10 s
Training iterations	3,000	3,000	3,000

Note: Cassie’s Baseline DR already includes mass and gain randomization (unlike Go2 Baseline which has friction-only), reflecting the higher inherent difficulty of bipedal transfer.

3.3.6 Observation Noise (Training Only)

During training, sensor noise is injected to improve robustness:

Table 43: Cassie Observation Noise Configuration

Sensor	Noise Level
Linear velocity	± 0.1 m/s
Angular velocity	± 0.2 rad/s
Projected gravity	± 0.05
Joint positions	± 0.01 rad
Joint velocities	± 1.5 rad/s

3.3.7 MuJoCo Model Details

The MuJoCo Cassie model (`external/agility_cassie/cassie.xml`) has several features critical to the transfer:

Timestep: $\Delta t = 0.0005$ s (2 kHz), required for constraint stability. This is $10\times$ faster than Go2’s 0.005 s timestep.

4-Bar Linkage Constraints: Implemented via four `<connect>` equality constraints linking plantar rods to foot bodies and achilles rods to heel springs. The solver parameter `solref="0.005 1"` defines the constraint resolution timescale.

Motor Actuators: Cassie uses MuJoCo `<motor>` actuators where the control signal is related to torque by the gear ratio:

$$F = \text{gear} \times \text{clip}(\text{ctrl}, \text{ctrlrange}) \quad (6)$$

The gear ratios vary by joint (hip-roll: 25, hip-pitch/knee: 16, foot: 50), creating non-uniform torque limits:

Table 44: Cassie MuJoCo Motor Actuator Configuration

Joint	Gear	ctrlrange	Max Torque	Isaac Limit
hip-roll	25	± 4.5	112.5 N·m	200 N·m
hip-yaw	25	± 4.5	112.5 N·m	200 N·m
hip-pitch	16	± 12.2	195.2 N·m	200 N·m
knee	16	± 12.2	195.2 N·m	200 N·m
foot	50	± 0.9	45.0 N·m	20 N·m

The hip-roll torque limit mismatch (112.5 vs. 200 N·m) is particularly problematic for lateral balance—the robot can lose the ability to prevent sideways tipping.

Passive Structures: The model includes leaf springs (`left/right-shin` with stiffness 1500 N·m/rad) and heel springs (stiffness 1250 N·m/rad) that store energy during locomotion. These springs are not present in the Isaac Lab model but contribute to Cassie’s gait dynamics in MuJoCo.

3.3.8 Observation Construction in MuJoCo

The Cassie MuJoCo deployment script constructs observations differently from Go2 in several key respects:

1. **Linear velocity:** Read from `qvel[0:3]` (world frame), rotated to body frame via `quat_rotate_inverse` with the corrected minus-sign formula. **No scaling applied** (unlike Go2’s $\times 2.0$).
2. **Angular velocity:** Read from `qvel[3:6]` (world frame), rotated to body frame. **No scaling applied** (unlike Go2’s $\times 0.25$).
3. **Joint positions:** Offset from MuJoCo home keyframe angles (not Isaac Lab defaults). For passive ankle joints, the tarsus joint position is used as a proxy, since tarsus and ankle are kinematically linked through the 4-bar mechanism.
4. **Joint velocities:** Direct readout with no scaling (unlike Go2’s $\times 0.05$).
5. **Previous actions:** Stored from the last policy output. In the 12-action variant, ankle actions (indices 4, 10) are set to zero since they cannot be applied. In the 10-action passive variant, all 10 actions are meaningful.
6. **Safety clip:** All observations are clipped to ± 100 to prevent transient numerical spikes from destabilizing the policy. This clip is not applied during training but prevents rare MuJoCo integration artifacts from corrupting the observation.

3.3.9 PD Control Implementation

The Cassie MuJoCo PD control loop differs from Go2 in both gain structure and torque application:

```
# Compute PD torque (with 6x gain scale for implicit->explicit PD)
target = mj_default + action_scale * action
torque = Kp_scaled * (target - q_current) - Kd_scaled * q_vel
torque = clip(torque, -effort_limit, effort_limit)

# Convert to MuJoCo motor control signal
ctrl = torque / gear_ratio
ctrl = clip(ctrl, ctrlrange_min, ctrlrange_max)
```

Listing 3: Cassie PD control with gain scaling and motor model

The `pd_gain_scale=6.0` factor compensates for the difference between Isaac Lab’s implicit PD (constraint-based, very stiff) and MuJoCo’s explicit PD (pre-step torque, compliant). Without this scaling, the MuJoCo robot behaves as if its joints are “loose,” unable to hold positions against gravity and dynamic loads. This $6\times$ factor was determined empirically by matching the step response of individual joints between simulators.

3.4 Transfer Debugging Journey: Four Critical Bugs

The Cassie transfer required identifying and fixing four independent bugs, each of which was investigated systematically through controlled experiments.

3.4.1 Bug 1: *quat_rotate_inverse* Sign Error

Discovery: The initial deployment used a `quat_rotate_inverse` function with an incorrect sign in the formula. The function computed the *forward* rotation ($v + w \cdot t$) instead of the *inverse* ($v - w \cdot t$), where w is the scalar part of the quaternion and t is the cross-product term.

Verification: With the robot pitched 30° forward and `qvel[3 : 6] = [0, 0, 1]`:

- Wrong (forward rotation): `ang_vel_b = [-0.5, 0, 0.866]` (distorted)
- Correct (inverse rotation): `ang_vel_b = [0, 0, 1]` (unchanged, since `qvel[3:6]` is already in body frame for MuJoCo free joints)

Impact: This bug affected the body-frame linear velocity observation, corrupting 3 of 48 observation dimensions.

3.4.2 Bug 2: *Angular Velocity Frame Convention*

Discovery: The deployment code rotated `qvel[3:6]` (angular velocity) from “world frame” to body frame. However, MuJoCo’s free joint convention stores angular velocity in *world frame* for `qvel[3:6]`, requiring rotation to body frame—but the Go2 deployment (which worked correctly) also rotated it.

Investigation: Careful analysis revealed that for MuJoCo free joints:

- `qvel[0:3]` = linear velocity in **world frame** → needs `quat_rotate_inverse` to body frame
- `qvel[3:6]` = angular velocity in **world frame** → needs `quat_rotate_inverse` to body frame

The bug was actually in combination with Bug 1: the *forward* rotation was being applied instead of the *inverse*, which for angular velocity produced a doubly-wrong result.

Fix: Correct the sign in `quat_rotate_inverse` (Bug 1 fix) and apply it to both linear and angular velocity.

3.4.3 Bug 3: *Observation Baseline (Isaac Lab Defaults vs. MuJoCo Defaults)*

Discovery: The deployment script computed joint position observations as:

$$q_{\text{obs}} = q_{\text{current}} - q_{\text{isaac_default}}$$

But MuJoCo starts from its own home keyframe, which differs by up to 0.6 rad from Isaac Lab defaults (Table 36).

Experiment: Switching to MuJoCo home defaults:

$$q_{\text{obs}} = q_{\text{current}} - q_{\text{mj_default}}$$

Result: Robot survival improved from **16 steps to 153 steps** and achieved forward velocity up to 2.4 m/s before becoming unstable at high speed.

Interpretation: This was the single most impactful fix. The 0.5–0.6 rad offset in the observation baseline meant the policy “saw” a fundamentally different pose from what it expected, causing it to apply corrective actions for a postural error that did not physically exist.

3.4.4 Bug 4: Timestep Too Large (Constraint Solver Instability)

Discovery: Using $\Delta t = 0.005$ s (matching Go2 and the Isaac Lab training timestep) caused Cassie’s 4-bar linkage constraints to violate within the first few steps.

Root cause: The Cassie MuJoCo XML specifies `solref="0.005 1"` for equality constraints. At $\Delta t = 0.005$ s, the constraint solver has only $0.005/0.005 = 1$ iteration to resolve—far too few for a kinematic loop constraint. At the native $\Delta t = 0.0005$ s, the solver gets $0.005/0.0005 = 10$ iterations, sufficient for convergence.

Fix: Keep MuJoCo’s native timestep ($\Delta t = 0.0005$ s) and increase the control decimation from 4 to 40 to maintain the 50 Hz policy rate.

3.4.5 Summary of Bug Impact

Table 45: Cassie Transfer Bug Summary

Bug	Description	Before Fix	After Fix
1	quat_rotate_inverse sign	Immediate collapse	—
2	Angular velocity frame convention	Corrupted observations	—
3	Observation baseline (Isaac vs. MuJoCo defaults)	16 steps	153 steps
4	Timestep $\rightarrow$ constraint solver coupling	Linkage explosion	Stable constraints

3.5 Experiment 2: Cassie Transfer Analysis

3.5.1 Hypotheses

Hypothesis 13 (Structural Mismatch Dominance). *The Cassie sim-to-sim gap will be larger than Go2’s, driven by the structural mismatch (12 $\rightarrow$ 10 joints, different defaults) rather than actuator dynamics alone.*

Hypothesis 14 (Observation Baseline Sensitivity). *Using MuJoCo home keyframe angles (rather than Isaac Lab defaults) as the observation baseline will significantly improve transfer performance, because the policy’s joint position observations must reflect the actual kinematic state relative to a physically meaningful reference pose.*

Hypothesis 15 (4-Bar Linkage as Fundamental Barrier). *The passive ankle joints, constrained by the 4-bar linkage, will create a fundamental transfer barrier that cannot be fully resolved by domain randomization, because the policy learns to use all 12 action dimensions during training but can only actuate 10 during deployment.*

3.5.2 Experiment Setup

Three controlled experiments were conducted, each changing a single variable:

Table 46: Cassie Transfer Experiments

Exp	Change	Description	Metric
1	Fix angular velocity bug	Remove incorrect rotation of qvel[3:6]	Survival steps
2	Use MuJoCo defaults	Offset observations from MuJoCo home pose instead of Isaac Lab defaults	Survival steps, max velocity
3	Correct joint mapping	Fix joint index mapping to verified values	Survival steps

All experiments used the Baseline policy (3,000 iterations) with forward velocity command ($v_x = 0.5$ m/s, $v_y = 0$, $\omega_z = 0$).

3.5.3 Results

Table 47: Cassie Transfer Experiment Results

Exp	Configuration	Survival	Max v_x	Outcome
—	All bugs present	<5 steps	N/A	Immediate collapse
1	+ Ang. vel. fix	16 steps	N/A	Still fails quickly
2	+ MuJoCo defaults	153 steps	2.4 m/s	Walks, unstable at high speed
3	+ Correct joint mapping	22 steps	~0.5 m/s	Policy struggles with dynamics

3.5.4 Analysis

Structural Mismatch—Supported. The Cassie transfer gap is qualitatively larger than Go2’s. While Go2 achieved 100% survival once the pipeline bugs were fixed, Cassie achieves only partial success (153 steps maximum) even with all known bugs fixed. This is consistent with the structural mismatch (12 $\rightarrow$ 10 joints, different defaults) creating a more fundamental gap than Go2’s actuator-only mismatch.

Observation Baseline—Strongly supported. Switching from Isaac Lab defaults to MuJoCo home keyframe angles improved survival by **9.6 $\times$** (16 $\rightarrow$ 153 steps) and enabled forward locomotion up to 2.4 m/s. This is the single most impactful fix in the entire Cassie pipeline. The observation baseline effectively defines “what the policy thinks zero looks like”—using the wrong baseline is equivalent to adding a constant bias to all joint position observations.

4-Bar Linkage Barrier—Partially supported. The passive ankle joints do create a fundamental limitation: 2 of 12 action dimensions are wasted (16.7% of the action space). However, the 153-step survival with forward velocity up to 2.4 m/s suggests that the policy can *partially* adapt—the ankle actions are small enough in practice that dropping them does not immediately destabilize the gait. The instability at high speed may be related to the ankle constraint, but further investigation is needed to confirm.

Remaining challenges and additional bugs identified. Even after the first four bugs were fixed, further investigation revealed additional mismatch sources that prevented 100% transfer success:

Table 48: Additional Cassie Transfer Issues Identified After Initial Bug Fixes

Bug	Issue	Description and Fix
5	Action clip too permissive	Policy actions clipped at ± 5 allowed divergence via feedback loop in observation indices 36–47 (previous actions). Fixed to ± 1.0 .
6	Implicit vs. Explicit PD gain mismatch	Isaac Lab uses constraint-based PD (very stiff, implicit). MuJoCo uses external PD (needs higher gains). Fix: apply <code>pd_gain_scale=6.0</code> to all gains.
7	Motor <code>ctrlrange</code> too small	Cassie XML hip-roll limited to 112.5 N·m ($\text{gear} \times \text{range} = 25 \times 4.5$) but Isaac Lab allows 200 N·m. Fix: expand <code>ctrlrange = effort_limit / gear</code> .
8	Passive ankle mismatch (fundamental)	Isaac Lab actuates ankles (12 actions); MuJoCo has passive 4-bar linkage (10 actuators). The policy wastes 2 action dimensions on joints it cannot control.

3.6 Proposed Solution: Passive Ankle Training

The most fundamental barrier to Cassie transfer is Bug 8: the policy is trained with 12 actuated joints in Isaac Lab, but only 10 are controllable in MuJoCo. Rather than dropping ankle actions at deployment and hoping the policy compensates, the principled solution is to *train with passive ankles from the start*, so that the policy never learns to rely on ankle actuation.

3.6.1 Implementation

A passive ankle training variant was fully implemented with two new components:

1. Environment configuration (CassieSim2SimPassiveEnvCfg):

- Removes ankle actuators from the Isaac Lab robot definition in `__post_init__`
- Reduces action space from 12 to **10 dimensions**
- Reduces observation space from 48 to **46 dimensions** (previous actions shrink from 12 to 10)
- Joint position and velocity observations remain 12-dimensional (passive ankle positions are still observed via the simulation)

2. Environment wrapper (CassiePassiveEnv):

- Maps 10-dimensional policy actions to 12-dimensional joint targets
- Automatically identifies active vs. passive joint indices by name (“ankle”)
- Uses `index_add_` for efficient batched assignment of actions to the correct joint slots
- Passive joints receive zero offset from their default positions (held at rest)

3. Deployment configuration (cassie_passive.yaml):

- Network input: 46 observations, output: 10 actions
- PD gains with `pd_gain_scale=6.0` to compensate for implicit/explicit PD mismatch

- Higher K_d values (hips: 10.0, knees: 20.0, toes: 5.0) vs. original (hips: 3.0, knees: 6.0, toes: 1.0) for additional damping
- Toe effort limit increased to 45.0 N·m (from 20.0) to better match MuJoCo’s foot actuator range

The environment is registered as `Isaac-Velocity-Flat-Cassie-Sim2Sim-Passive-v0`.

3.6.2 Expected Benefits

1. **Eliminates wasted action dimensions:** The policy uses all 10 output dimensions for joints it can actually control, improving sample efficiency and action precision.
2. **Observation consistency:** Both training (Isaac Lab) and deployment (MuJoCo) have passive ankles, eliminating the observation mismatch for previous actions (10-dim in both, not 12-dim training vs. 10-dim deployed).
3. **Learns ankle-free balance:** The policy must learn to balance and walk without ankle torque, developing strategies that rely on hip, knee, and toe actuation—the exact set available in MuJoCo.

Status: The passive training infrastructure is fully implemented and registered. Training has been completed, and the policy has been successfully deployed and tuned in MuJoCo.

3.6.3 Passive Policy Tuning and Standing Success

After training the passive ankle policy in Isaac Lab, extensive deployment tuning was performed to achieve stable standing in MuJoCo. This section documents the debugging journey and the critical fixes that enabled the policy to maintain balance for extended periods.

Critical Bug: Observation Reference Mismatch The most impactful discovery was a subtle but critical mismatch in observation computation. The policy’s joint position observations are computed as:

$$\text{obs_joint_pos}_i = q_{\text{current},i} - q_{\text{reference},i} \quad (7)$$

Initially, the deployment script used Isaac Lab’s default joint angles as the reference (e.g., `thigh_joint` = -1.8 rad for the crouched training pose). However, MuJoCo initializes the robot from its “home” keyframe, which has significantly different angles (e.g., `thigh_joint` = -1.2 rad for a more upright pose).

The fix: Use MuJoCo’s home keyframe angles as the observation reference, not Isaac Lab’s defaults:

Table 49: Critical Joint Reference Differences

Joint	Isaac Lab Ref	MuJoCo Home	Difference
hip_flexion	+1.0 rad	+0.497 rad	−0.503 rad
thigh_joint (knee)	−1.8 rad	−1.2 rad	+0.6 rad
ankle_joint	+1.57 rad	+1.427 rad	−0.143 rad

With the old (incorrect) reference, the policy received observations indicating the robot was 0.6 rad “too straight” at the knees, causing it to command a deep crouch that led to collapse within 0.7–0.9 seconds. With the corrected reference, the policy correctly interprets its starting pose and maintains balance.

Parameter Tuning Results Systematic parameter tuning was performed to optimize stability:

Table 50: Passive Policy Tuning Summary

Parameter	Optimal Value	Effect
pd_gain_scale	8.0	Balance between stiffness and stability
action_scale	0.3	Reduced from training (0.5) to dampen response
toe_joint gains	[40.0, 15.0, 100.0]	Increased from [20, 10, 100] for foot control
Ankle observations	Zeroed	Tarsus feedback made stability worse
Physics frequency	2000 Hz	Required for 4-bar linkage constraint stability

Standing Performance With the corrected reference angles and optimized parameters, the passive Cassie policy achieved the following results:

- **Survival time:** 1.70 seconds (up from <1.0s with incorrect reference)
- **Height maintenance:** 0.926 m at $t = 1.0$ s (near home keyframe height of 1.0 m)
- **Failure mode:** Gradual forward drift (~ 0.3 – 0.6 m/s) leading to eventual collapse

Table 51: Passive Policy Performance Progression

Configuration	Survival Time	Height at $t = 1.0$ s
Incorrect reference (Isaac Lab defaults)	0.7–0.9 s	0.4–0.5 m (crouched)
Correct reference (MuJoCo home)	1.56 s	0.92 m
+ Optimized gains & toe PD	1.70 s	0.926 m

Feedforward Bias Experiments Several correction strategies were attempted to counteract the residual forward drift:

1. **Static feedforward bias:** Adding constant offsets to hip_flexion actions to induce backward lean. This reduced forward velocity but also reduced standing height, resulting in worse overall stability.
2. **Velocity-proportional correction:** A P-controller layer that adds backward lean proportional to forward velocity. This provided marginal improvement ($1.70 \rightarrow 1.72$ s) but did not fundamentally resolve the drift.
3. **Command manipulation:** Setting backward velocity commands ($v_x = -0.3$ m/s) paradoxically increased forward velocity, suggesting the policy interprets commands differently across the sim gap.

Key insight: The residual forward drift represents a fundamental sim-to-sim gap that cannot be fully bridged through deployment-time corrections. The most effective solutions would require retraining with either (1) more aggressive domain randomization in Isaac Lab, or (2) fine-tuning directly in MuJoCo.

Final Configuration The validated deployment configuration (`cassie_passive.yaml`) achieves consistent 1.70s standing:

```

pd_gain_scale: 8.0          # 8x gain scale for explicit PD
action_scale: 0.3           # Reduced from 0.5 (training)
simulation_dt: 0.001        # 1000 Hz physics timestep
control_decimation: 20      # 50 Hz policy rate
toe_joint: [40.0, 15.0, 100.0] # [Kp, Kd, effort_limit]
# Joint reference: MuJoCo home keyframe (not Isaac Lab defaults)

```

Listing 4: Key parameters for passive Cassie deployment

3.6.4 Model Investigation: Could a Different MuJoCo Model Improve Transfer?

Having achieved 1.70s standing with the passive ankle policy but facing persistent forward drift that deployment-time corrections could not resolve, I investigated whether the choice of MuJoCo model itself might be contributing to the sim-to-sim gap. This was motivated by the Go2 experience, where switching from `go2_description` to the official Unitree model improved transfer quality significantly.

Hypothesis Perhaps the `external/agility_cassie` model (from MuJoCo Menagerie) has simplifications or approximations that create unnecessary mismatch with Isaac Lab’s Cassie model. An alternative model—`cassie-mujoco-sim` from the Dynamic Robotics Laboratory at University of Michigan (Cassie’s original developers)—might better match Isaac Lab and enable longer standing times.

Model Comparison Two MuJoCo Cassie models were systematically compared:

1. `external/agility_cassie/cassie.xml` (MuJoCo Menagerie)
 - Source: Agility Robotics via DeepMind’s MuJoCo Menagerie
 - Base joint: `freejoint` (6 DOF)
 - Mesh format: OBJ files with textures
 - Default pose: "home" keyframe at $z = 1.0$ m
 - Collision: Multiple capsules per limb (detailed contact geometry)
2. `cassie-mujoco-sim/model/cassie.xml` (Dynamic Robotics Lab)
 - Source: University of Michigan (original Cassie simulator developers)
 - Base joint: $3 \times$ slide + $1 \times$ ball (equivalent 6 DOF)
 - Mesh format: STL files
 - Default pose: No keyframe defined ($z = 1.01$ m in body definition)
 - Collision: Single capsule per limb (simplified)

Table 52: Cassie MuJoCo Model Detailed Comparison

Property	cassie-mujoco-sim	external/agility_cassie	Identical?
Controllable joints	10 (5 per leg)	10 (5 per leg)	✓
Joint names	hip-roll, hip-yaw, etc.	hip-roll, hip-yaw, etc.	✓
Actuator gears	[25, 25, 16, 16, 50]	[25, 25, 16, 16, 50]	✓
Actuator ctrlrange	[± 4.5 , ± 12.2 , ± 0.9]	[± 4.5 , ± 12.2 , ± 0.9]	✓
4-bar linkage	solref=0.005 1	solref=0.005 1	✓
Pelvis mass	10.33 kg	10.33 kg	✓
Pelvis COM	[0.051, 0.0003, 0.028]	[0.051, 0.0003, 0.028]	✓
Shin spring	1500 N·m/rad	1500 N·m/rad	✓
Heel spring	1250 N·m/rad	1250 N·m/rad	✓
Timestep	dt=0.0005 s	dt=0.0005 s	✓
Default keyframe	None	"home" keyframe	×
Collision detail	Single capsule/limb	Multiple capsules/limb	×
Mesh format	STL	OBJ + textures	×

Key Finding: Models Are Functionally Identical The comparison revealed that both models implement **identical physics**:

- All 10 controllable joints with exact same names
- All actuator parameters match (gear ratios, control ranges, torque limits)
- All constraint parameters match (4-bar linkage solref values)
- All inertial properties match (mass, COM positions, inertia tensors)
- All passive spring stiffnesses match (shin 1500, heel 1250)
- Same required timestep (dt=0.0005 s for constraint stability)

The only differences are:

1. **Default keyframe:** `external/agility_cassie` provides a validated "home" pose; `cassie-mujoco-sim` does not (must manually initialize).
2. **Collision geometry:** `external/agility_cassie` uses multiple capsules per limb for more accurate ground contact; `cassie-mujoco-sim` uses simplified single capsules.
3. **Visual meshes:** Different file formats (OBJ vs. STL), no impact on physics.

Conclusion: Model Choice Is Not the Bottleneck Unlike Go2 (where model choice significantly affected transfer), for Cassie the model choice is **irrelevant to transfer quality** because:

1. **Both models are official:** MuJoCo Menagerie (DeepMind) and `cassie-mujoco-sim` (UMich) both validated against the real Cassie robot.
2. **Physics parameters are byte-for-byte identical:** All actuator, constraint, inertial, and spring properties match exactly—they were copied from the same original Agility Robotics specifications.
3. **The sim-to-sim gap is architectural, not parametric:** The 1.70s survival limitation stems from PhysX (Isaac Lab's implicit constraint-based PD) vs. MuJoCo PGS solver (explicit torque-based PD). This fundamental difference exists regardless of which XML file is used.
4. **Joint name remapping already handled:** The deployment script already maps Isaac Lab names (`hip_abduction`, `thigh_joint`, etc.) to MuJoCo names (`hip_roll`, `knee`, etc.). Both MuJoCo models use the same joint names, so switching models does not change this mapping.

5. **Collision geometry difference is minor:** While `external/agility_cassie` has more detailed collision (multiple capsules), both models produce qualitatively similar ground contact behavior. The forward drift problem occurs in both models.

Decision and Rationale Continue using `external/agility_cassie/cassie.xml` because:

- Has "home" keyframe for easier initialization
- More accurate collision geometry (better ground contact simulation)
- Already fully debugged and validated with current deployment pipeline
- Part of actively maintained MuJoCo Menagerie (regular updates from DeepMind)
- Slightly better visual quality (OBJ meshes with textures)

Switching to `cassie-mujoco-sim` would provide **zero physics benefit** while requiring additional manual initialization code to handle the missing keyframe.

Implications for Next Steps Since model choice cannot improve transfer performance, the 1.70s survival with forward drift represents a **fundamental sim-to-sim gap** that requires **training-time solutions**, not deployment-time fixes:

1. **Enhanced domain randomization:** Train with wider ranges of actuator gains, mass distributions, and joint friction to cover the PhysX–MuJoCo dynamics gap. Add explicit forward drift penalty to reward shaping.
2. **Simplified task:** Create a "standing only" task (commands always $[0, 0, 0]$) to remove gait complexity and focus purely on balance transfer.
3. **Direct MuJoCo fine-tuning:** As a last resort, fine-tune the policy in MuJoCo itself (though this defeats the purpose of sim-to-sim transfer research).

The next sections describe the implementation and results of strategies (1) and (2).

3.6.5 Enhanced Domain Randomization for Improved Transfer

Having ruled out model choice as a limiting factor, I turn to improving the training process itself. The baseline passive ankle policy achieved 1.70s standing, limited by:

1. Forward drift (0.3–0.6 m/s) despite zero velocity command
2. Eventual loss of balance and forward fall around 1.7s
3. Insufficient robustness to MuJoCo’s explicit PD dynamics

Hypothesis: More aggressive domain randomization during training can better cover the PhysX–MuJoCo sim-to-sim gap, and adding explicit penalties for the observed failure modes (forward drift) will guide the policy toward transfer-robust behaviors.

Enhanced Training Configuration An enhanced passive ankle training variant (`Isaac-Velocity-Flat-Cassie-Sim2Sim-Passive-Enhanced-v0`) was created with:

Table 53: Enhanced Domain Randomization vs. Baseline

Parameter	Baseline	Enhanced	Rationale
Actuator stiffness scale	[0.8, 1.2]	[0.5, 1.5]	Cover wider PD gain range
Actuator damping scale	[0.8, 1.2]	[0.5, 1.5]	Train for variable compliance
Leg mass variation	None	± 0.5 kg	COM shift randomization
Joint friction	Fixed	[0.0, 2.0]	Variable energy dissipation
Push frequency	10–15s	5–10s	2× more disturbances
Forward drift penalty	0.0	−1.0	Penalize v_x when cmd=0
Standing bonus	0.0	+0.5	Reward stillness

New Reward Terms Two reward terms specifically target the observed MuJoCo failure mode:

1. **Forward drift penalty** (weight = −1.0):

$$r_{\text{drift}} = -1.0 \cdot \begin{cases} v_x^2 & \text{if } |v_{x,\text{cmd}}| < 0.1 \\ 0 & \text{otherwise} \end{cases} \quad (8)$$

Penalizes forward velocity specifically when commanded to stand still, directly addressing the drift failure mode.

2. **Standing success bonus** (weight = +0.5):

$$r_{\text{stand}} = +0.5 \cdot \mathbb{I}_{|v_{\text{cmd}}| < 0.1 \text{ and } |v_{\text{body}}| < 0.1} \quad (9)$$

Rewards achieving stillness when commanded to stand, encouraging the policy to learn zero-velocity balance.

Status Training completed for 30,000 iterations (8–10 hours). Deployment testing pending. Expected improvements: survival >3.0s, drift <0.2 m/s.

3.6.6 Simplified Task: Balance Standing Only

The second training-time solution is to simplify the task itself. **Hypothesis:** A policy trained exclusively to stand still (no locomotion) will transfer more easily because:

1. Smaller action magnitudes → less sensitivity to PD gain mismatch
2. No gait dynamics → no compound errors across steps
3. Clear success criterion: $|v| \approx 0$
4. Independent timesteps: each control decision is stateless (no gait phase tracking)

Task Definition A dedicated balance standing task (Isaac-Balance-Standing-Cassie-v0) was created:

- **Commands:** Always $[0, 0, 0]$ (no resampling)
- **Episode length:** 20 seconds
- **Reward structure:**
 - Stillness bonus: +2.0 when $\sqrt{v_x^2 + v_y^2} < 0.05$ m/s
 - Linear velocity penalty: $-5.0 \cdot (v_x^2 + v_y^2)$
 - Angular velocity penalty: $-0.2 \cdot (\omega_x^2 + \omega_y^2 + \omega_z^2)$
 - Orientation penalty: $-8.0 \cdot \|g_{\text{proj}} - [0, 0, -1]\|^2$
 - Action smoothness: $-0.05 \cdot \|a_t - a_{t-1}\|^2$

- **Network:** Same as locomotion (256–256–256, ELU activation)
- **Action/Obs space:** 10 actions, 46 observations (passive ankles)

Table 54: Standing Task vs. Locomotion Task

Aspect	Locomotion	Standing Only
Command range	$v_x \in [-0.8, 0.8]$ m/s	$v_x = 0$ (fixed)
Task complexity	Gait generation	Balance only
Action magnitudes	Large (propulsion)	Small (corrections)
PD gain sensitivity	High	Low
Compound errors	Yes (gait phase drift)	No (stateless)
Training time	30K iter (8–10h)	20K iter (5–6h)
<i>MuJoCo transfer hypothesis:</i>		
Expected survival	1.70s (measured)	>10s (target)

Status and Expected Outcomes Training in progress. If the standing-only policy achieves >10s survival in MuJoCo, it would:

1. **Validate the hypothesis** that task complexity (not fundamental physics) limits transfer
2. **Enable curriculum learning:** Standing (0,0,0) $\rightarrow$ Slow forward (0–0.3 m/s) $\rightarrow$ Full velocity (0–0.8 m/s)
3. **Provide a baseline** for measuring the impact of enhanced domain randomization

3.7 Cassie Transfer: Complete Progression and Lessons

The Cassie sim-to-sim transfer required a systematic multi-stage progression. Table 55 summarizes the complete journey from immediate collapse to training-time solutions:

Table 55: Cassie Transfer: Complete Progression from Failure to Improved Training

Stage	Intervention	Survival	Factor	Key Insight Gained
0	Baseline (all bugs)	<5 steps	—	Immediate collapse
<i>Phase 1: Deployment Bug Fixes</i>				
1	Fix quat_rotate sign	16 steps	3.2×	Obs corruption critical
2	MuJoCo home defaults	153 steps	9.6×	Baseline is most impactful
3	Timestep dt=0.0005s	Stable 4-bar	Qual.	Constraints need fine dt
4	Expand motor ctrlrange	No saturation	Qual.	Torque limits matter
<i>Result: 153 steps with 12 actions, but unstable at 2.4 m/s</i>				
<i>Phase 2: Structural Realignment (Retraining)</i>				
5	Train with passive ankles (10 actions, 46 obs)	1.70s 0.926m height	Structural match	Eliminate wasted actions But forward drift persists
<i>Result: Standing works, but 0.3–0.6 m/s forward drift remains</i>				
<i>Phase 3: Model Investigation</i>				
6	Compare cassie-mujoco-sim vs. agility_cassie	No change (as predicted)	0×	Models are identical Physics gap is architectural
<i>Conclusion: Need training-time solutions, not deployment fixes</i>				
<i>Phase 4: Training-Time Solutions (In Progress)</i>				
7	Enhanced DR (wider gains, drift penalty)	In progress Expected >3s	TBD	Target PhysX/MuJoCo gap via randomization
8	Standing-only task (cmd = [0,0,0] only)	In progress Expected >10s	TBD	Simplify before walking Enable curriculum

3.7.1 Process Flow: From Debugging to Redesign

The Cassie transfer exhibits a clear progression pattern:

Phase 1: Deployment bug hunting (Stages 0–4)

- Each bug required hypothesis-driven testing
- Observation baseline (Stage 2) had the largest single impact: 9.6× improvement
- Achieved partial success: 153 steps, 2.4 m/s forward velocity
- *Limitation revealed:* Policy cannot fully compensate for 12 → 10 action mismatch, unstable at high speed

Phase 2: Structural realignment via retraining (Stage 5)

- *Recognition:* Dropping 2 actions at deployment is fundamentally flawed
- *Solution:* Train with passive ankles from the start (10 actions, 46 obs)
- *Result:* 1.70s standing at 0.926m height, policy works as designed
- *New limitation revealed:* Forward drift (0.3–0.6 m/s) persists despite zero velocity command
- *Contribution:* Proves passive ankle training is viable; isolates residual sim-to-sim gap

Phase 3: Model validation (Stage 6)

- *Question:* Could model choice explain the remaining gap?
- *Method:* Systematic comparison of two official Cassie models (cassie-mujoco-sim vs. agility_cassie)

- *Result*: Models are byte-for-byte identical in all physics parameters
- *Contribution*: Rules out model choice as a variable; confirms gap is PhysX vs. MuJoCo solver architecture
- *Contrast with Go2*: For Go2, model choice mattered significantly; for Cassie, it does not

Phase 4: Training-time solutions (Stages 7–8, ongoing)

- *Recognition*: 1.70s with drift cannot be fixed post-hoc
- *Strategy A* (Enhanced DR): Wider actuator gain range, leg mass variation, explicit drift penalty
- *Strategy B* (Simplified task): Standing-only (no locomotion) to reduce action magnitudes and eliminate gait complexity
- *Expected contribution*: If standing achieves >10s, enables curriculum learning (stand → slow walk → full locomotion)

3.7.2 Key Insights: Co-Design of Training and Deployment

The Cassie experience demonstrates that successful sim-to-sim transfer requires **iterative co-design**:

1. **Deployment bugs must be fixed first** (Stages 1–4): Without correct observations and physics, no amount of training will succeed.
2. **Training must match deployment constraints** (Stage 5): The 12 → 10 action mismatch was fundamental; training with passive ankles from the start was necessary, not optional.
3. **Model choice must be validated** (Stage 6): Unlike Go2, Cassie’s model choice is irrelevant because both official models are identical. This rules out "try a different model" as a solution.
4. **Domain randomization must target the specific gap** (Stage 7): Generic DR is insufficient; must explicitly randomize actuator gains to cover PhysX vs. MuJoCo PD dynamics, and add failure-mode-specific penalties (forward drift).
5. **Task complexity may need reduction** (Stage 8): For difficult transfers, starting with a simpler task (standing) before attempting the full task (locomotion) may be necessary.

This iterative cycle—debug → retrain → validate → redesign training—is fundamentally different from Go2, where all issues were deployment-time fixes. Cassie’s structural complexity (passive ankles, 4-bar linkage, higher DOF, kinematic constraints) makes it a more challenging but more representative case study for sim-to-sim transfer research.

3.7.3 Summary of Findings

1. Cassie transfer is **qualitatively harder** than Go2: 8+ bugs vs. 3 bugs, requires retraining vs. deployment fixes only.
2. The **observation reference baseline** (MuJoCo home vs. Isaac Lab defaults) was the most impactful single fix: 9.6× improvement.
3. The **4-bar linkage** creates dual challenges: passive joints (lost actuator authority) and constraint-timestep coupling (requires $dt=0.0005s$).
4. The **implicit vs. explicit PD gap** is severe: requires 8× gain scaling to approximate Isaac Lab’s constraint-based PD.
5. **Passive ankle retraining** makes progress on structural mismatch, achieving 1.70s of standing (limited by residual forward drift).
6. **Model choice is irrelevant**: Both official models (cassie-mujoco-sim and agility_cassie) are byte-for-byte identical in physics.

7. **Residual forward drift** ($\sim 0.3\text{--}0.6$ m/s) represents an irreducible sim-to-sim gap requiring training-time solutions (enhanced DR, simplified tasks).
8. Dominant mismatch source is **kinematic structure** (passive joints, pose conventions, constraint physics), not actuator dynamics alone.

4 Part 3: Flamingo Two-Wheel-Legged Robot

4.1 Robot Overview

The Flamingo is a two-wheel-legged robot—a dynamically unstable platform that must actively balance while locomoting, combining characteristics of wheeled robots (continuous high-speed motion) and legged robots (variable height, terrain adaptation). Unlike the Go2 and Cassie, the Flamingo uses a **hybrid control architecture**: position control for hip, shoulder, and leg joints, and *velocity control* for the wheel joints.

Table 56: Flamingo Robot Specifications

Property	Value
Type	Two-wheel-legged robot (dynamically unstable)
DOF	8 (per side: hip, shoulder, leg, wheel)
Control modes	Position (hip, shoulder, leg) + Velocity (wheel)
Leg gear ratio	-1.5 (mechanical transmission)
Standing height target	0.45 m
Observation stacking	$3 \text{ frames} \times 28 \text{ dims} + 4 \text{ cmd} = 88 \text{ total}$
Action scale	1.0 (position), 20.0 (wheel velocity)

Why is Flamingo the hardest transfer? Three compounding challenges make this the most difficult of the three robots:

1. **Dynamic instability:** The robot must actively balance on two wheels—any observation or control mismatch causes immediate fall-over.
2. **Heterogeneous actuators:** The policy must simultaneously produce position targets (joints) and velocity targets (wheels) through a single output vector.
3. **Gear ratio and axis conventions:** The leg joints have a -1.5 gear ratio that transforms PD control. Right-side joints in the MuJoCo XML have `axis="0 0 -1"`, which MuJoCo handles automatically (no manual sign flip needed for joint angles/velocities). However, wheels have an additional USD-to-XML axis mismatch requiring explicit negation (see Bug 4).

4.2 Transfer Challenges Unique to Flamingo

4.2.1 Challenge 1: Observation Stacking

Unlike Go2 and Cassie (which use single-frame observations), Flamingo’s policy uses **3-frame temporal stacking**: the current observation frame (28 dims) is concatenated with the two most recent frames, plus a 4-dimensional command vector, for a total of $3 \times 28 + 4 = 88$ input dimensions.

This requires the deployment script to maintain a rolling buffer of observation history, with the *newest frame first* in the concatenation order—a subtle ordering requirement that, if reversed, completely breaks the policy’s temporal reasoning.

4.2.2 Challenge 2: Gear Ratio (-1.5) for Leg Joints

The Flamingo’s leg joints use a mechanical transmission with gear ratio $g = -1.5$. This affects both observations and control:

- **Observations:** Joint positions and velocities must be multiplied by g before being fed to the policy (the policy was trained with gear-transformed values).
- **PD Control:** The effective PD gains are transformed by g^2 : $K_{p,\text{eff}} = g^2 K_p = 2.25 K_p$, $K_{d,\text{eff}} = g^2 K_d = 2.25 K_d$.

4.2.3 Challenge 3: Joint Axis Inversions

The MuJoCo XML defines several right-side joints with inverted axes:

- `right_shoulder_joint`: `axis="0 0 -1"` (inverted)
- `right_leg_joint`: `axis="0 0 -1"` (inverted)
- `right_wheel_joint`: `axis="0 0 -1"` (inverted)

These inversions require negating both the observation values (joint positions/velocities read from MuJoCo) and the control targets (actions sent to MuJoCo) for the right side. Missing a single negation creates asymmetric behavior where the robot curves or falls to one side.

4.2.4 Challenge 4: Hybrid Position-Velocity Control

The 8-dimensional action vector has two distinct control modes:

Table 57: Flamingo Action Space (8 dimensions)

Index	Joint	Mode	Scale	PD Gains
0	left_hip	Position	1.0	$K_p = 100, K_d = 1.5$
1	right_hip	Position	1.0	$K_p = 100, K_d = 1.5$
2	left_shoulder	Position	1.0	$K_p = 100, K_d = 1.5$
3	right_shoulder	Position	1.0	$K_p = 100, K_d = 1.5$
4	left_leg	Position	1.0	$K_p = 120, K_d = 1.5$ ($g = -1.5$)
5	right_leg	Position	1.0	$K_p = 120, K_d = 1.5$ ($g = -1.5$)
6	left_wheel	Velocity	20.0	$K_p = 0, K_d = 0.7$
7	right_wheel	Velocity	20.0	$K_p = 0, K_d = 0.7$

The wheel joints use *pure velocity control* ($K_p = 0$), meaning the control signal is a damping-based torque: $\tau_{\text{wheel}} = K_d(\dot{q}_{\text{target}} - \dot{q}_{\text{current}})$, where $\dot{q}_{\text{target}} = 20.0 \times a_{\text{wheel}}$. This is fundamentally different from the position control used for all Go2 and Cassie joints.

4.2.5 Challenge 5: Delayed PD Actuator

Isaac Lab’s Flamingo training uses a `DelayedPDActuator` that introduces random 0–4 step delays in the control signal, simulating real actuator latency. Direct PD control in MuJoCo has zero delay, creating an actuator dynamics mismatch that goes beyond the implicit/explicit PD difference seen with Go2.

4.3 Training Configuration

Table 58: Flamingo Training Configuration

Parameter	Value
Task	Flamingo_Flat_Stand_Drive (velocity tracking)
Physics timestep	0.005 s (200 Hz)
Num environments	4,096
Network architecture	Actor: [512, 256, 128], ELU
Input dimension	88 (3 frames $\times$ 28 + 4 cmd)
Output dimension	8 (6 position + 2 velocity targets)
Training iterations	5,000
Command range	v_x : $[-1.5, 1.5]$, v_y : 0.0 (disabled), ω_z : $[-2.5, 2.5]$

4.3.1 Observation Structure (28 dims per frame)

Table 59: Flamingo Per-Frame Observation (28 dimensions)

Index	Name	Dim	Scale	Note
0–3	hip_shoulder_joint_pos	4	$\times 1.0$	L/R hip, L/R shoulder
4–5	leg_joint_pos	2	$\times g$	$g = -1.5$ gear ratio
6–9	hip_shoulder_joint_vel	4	$\times 0.15$	Velocity scaling
10–11	leg_joint_vel	2	$\times g \times 0.15$	Gear + velocity scale
12–13	wheel_joint_vel	2	$\times 0.15$	Wheel angular velocity
14–16	base_ang_vel	3	$\times 0.25$	Body frame
17–19	projected_gravity	3	$\times 1.0$	Gravity in body frame
20–27	last_action	8	$\times 1.0$	Previous policy output

4.3.2 Reward Function

The Flamingo reward function emphasizes balance and height maintenance, reflecting the robot’s dynamic instability:

Table 60: Flamingo Reward Function (Flat_Stand_Drive)

Term	Weight	Purpose
track_lin_vel_xy_exp	+2.0	Exponential velocity tracking ($\sigma = 0.25$)
track_ang_vel_z_exp	+1.0	Exponential yaw rate tracking ($\sigma = 0.25$)
alive_bonus	+2.0	Reward per step for not falling
flat_orientation	-5.0	Keep robot level (roll + pitch)
base_height	-25.0	Target height 0.358 m
lin_vel_z_l2	-1.0	Penalize vertical bouncing
ang_vel_xy_l2	-0.05	Penalize roll/pitch rate
joint_deviation_hip	-1.0	Keep hips near zero
joint_deviation_shoulder	-0.5	Keep shoulders near zero
joint_deviation_leg	-0.5	Keep legs near zero
undesired_contacts	-0.5	Penalize body/leg ground contact
dof_torques_l2	-5.0×10^{-5}	Minimize energy consumption
dof_acc_l2	-2.5×10^{-7}	Smooth joint acceleration
action_rate_l2	-0.01	Smooth action changes
dof_pos_limits	-1.0	Avoid joint limits
termination_penalty	-200.0	Penalty for falling

The **base_height** weight (-25.0) is the second-largest penalty, reflecting how critical height control is for a wheeled-legged platform. The **alive_bonus** (+2.0 per step $\times$ 400 steps = 800 points potential) strongly incentivizes the policy to prioritize survival.

4.3.3 Domain Randomization

The Flamingo training includes:

- Static friction: [0.3, 1.0], dynamic friction: [0.3, 0.8]
- Base mass offset: ± 3.0 kg
- COM position offset: ± 0.02 – 0.05 m
- Actuator gain scaling: $\times [0.7, 1.3]$
- Actuator delay: 0–4 steps (via `DelayedPDActuator`)
- Reset noise: $\pm 25^\circ$ pitch/roll initial orientation

4.3.4 CO-RL Observation Stacking Mechanism

The Flamingo uses the CO-RL (Concurrent Online Reinforcement Learning) framework with temporal observation stacking. This is fundamentally different from Go2 and Cassie’s single-frame approach and introduces several transfer-critical details:

Stacking protocol: The CLI default `-num_policy_stacks 2` specifies 2 *additional* frames beyond the current frame, producing 3 total frames. These are concatenated in **reversed order** (newest first):

```
# History buffer: deque of 3 most recent observations
obs_history = deque(maxlen=3) # [oldest, middle, newest]

# Stacking for policy input: REVERSED
stacked = concatenate(list(reversed(obs_history)))
# Result: [newest(28), middle(28), oldest(28), command(4)] = 88 dims
```

Listing 5: CO-RL observation stacking (newest first)

Initialization: At episode reset, all 3 history slots are filled with the same initial observation. This ensures the policy does not see zeros or garbage data in early steps.

Why reversed order matters: The policy’s first 28 input dimensions correspond to the most recent observations, giving the network direct access to current state without requiring temporal reasoning. If the order is wrong (oldest first), the policy “sees” stale data as if it were current, producing catastrophically wrong actions.

4.3.5 MuJoCo Model and Axis Inversions

The MuJoCo Flamingo model (`flamingo_torque.xml`) defines several right-side joints with inverted axes:

Table 61: Flamingo Joint Axis Configuration

Joint	Axis	Inverted?
left_hip_joint	0 0 1	No
left_shoulder_joint	0 0 1	No
left_leg_joint	0 0 1	No
left_wheel_joint	0 0 1	No
right_hip_joint	0 0 1	No
right_shoulder_joint	0 0 -1	Yes
right_leg_joint	0 0 -1	Yes
right_wheel_joint	0 0 -1	Yes

A critical discovery was that **no manual sign negation is needed** in the deployment code. The inverted axes in the XML are a design choice that ensures symmetric actions produce symmetric motion. Positive control input on both legs produces outward/matching motion because the axis inversion is already encoded in the model. Early transfer attempts that manually negated right-side observations/actions produced asymmetric behavior (robot curving to one side).

Torque limit mismatch: The MuJoCo XML specifies a default `ctrlrange` of ± 23 N·m, but Isaac Lab training allows 60 N·m (joints) and 36 N·m (wheels). This must be corrected at simulation initialization by overriding:

```
model.actuator_ctrlrange[:6] = [[-60, 60]] # Joint actuators
model.actuator_ctrlrange[6:] = [[-36, 36]] # Wheel actuators
```

4.3.6 Deployment Control Flow

The Flamingo deployment loop runs at 50 Hz (policy rate) with $4\times$ physics substeps at 200 Hz:

1. **Build 28-dim observation** from MuJoCo state:
 - Hip/shoulder positions: direct readout (4 dims)
 - Leg positions: multiply by gear ratio $g = -1.5$ (2 dims)
 - Joint velocities: scale by 0.15 (6 dims), legs also by g
 - Wheel velocities: scale by 0.15 (2 dims)
 - Angular velocity: rotate to body frame, scale by 0.25 (3 dims)
 - Projected gravity: compute from quaternion (3 dims)
 - Previous action: 8-dim vector from last step
2. **Update history buffer:** Append new observation, oldest auto-discarded
3. **Stack observations:** Concatenate reversed buffer + command (88 dims)
4. **Policy inference:** Forward pass through actor network
5. **Clip actions:** ± 1.0 (matching training constraint)

6. **Apply PD control** (recomputed every physics substep):

- Joints (0–5): $\tau = K_p(q_{\text{target}} - q) - K_d\dot{q}$
- Legs (4–5): Use gear-transformed PD: $\tau_{\text{motor}} = K_p(a - q \cdot g) - K_d(\dot{q} \cdot g)$, then $\tau = \tau_{\text{motor}} \cdot g$
- Wheels (6–7): $\tau = K_d(\dot{q}_{\text{target}} - \dot{q})$ (velocity-only, $K_p = 0$)

7. **Step physics**: 4 substeps at $\Delta t = 0.005$ s

An important detail is that Isaac Lab recomputes PD torques **every physics substep** (inside the decimation loop), not just once per policy step. The MuJoCo deployment must replicate this behavior to match the effective control bandwidth.

4.4 Experiment 3: Flamingo Transfer Analysis

4.4.1 Hypotheses

Hypothesis 16 (Dynamic Instability Penalty). *The Flamingo will be the most sensitive to observation/control mismatches of all three robots, because any error in the balance-critical signals (angular velocity, gravity projection) immediately causes the dynamically unstable platform to fall over.*

Hypothesis 17 (Hybrid Control Challenge). *The hybrid position-velocity control architecture will create additional transfer difficulty compared to Go2 and Cassie (both pure position control), because the two control modes interact differently with MuJoCo’s physics solver.*

Hypothesis 18 (Height Target Unachievable). *The target standing height (0.45 m) will not be achieved in MuJoCo due to the compound effects of actuator dynamics mismatch, gear ratio imprecision, and the absence of the delayed actuator model—the robot will adopt a lower, more stable stance.*

4.4.2 Trial Process: Debugging Journey

CRITICAL LIMITATION: Before discussing the debugging process, it is essential to acknowledge that this experiment is fundamentally confounded by a **46% geometry mismatch** between the Isaac Lab training model (0.5562 m standing height) and all available MuJoCo XML files (0.2991 m standing height). This 0.256 m difference means the robot’s physical dimensions, link lengths, and mass distribution do not match between simulators. Analyzing subtle control details under this condition is logically questionable—the results reflect a compound interaction of geometry mismatch and control bugs, not a clean sim-to-sim transfer. The experiment proceeded to document the debugging methodology, but readers should interpret findings as illustrative of the debugging process rather than meaningful transfer validation.

The Flamingo transfer required identifying and fixing **8 critical bugs** that compounded to make the initial transfer completely non-functional. Unlike Go2 (3 bugs) and Cassie (8 bugs, but primarily kinematic), Flamingo’s bugs span observation construction, actuator dynamics, geometry mismatch, and control architecture—making it the most complex debugging process of the three robots.

Bug 1: Gyro Sensor Noise. The MuJoCo XML file `flamingo_torque.xml` defines a gyro sensor with `noise="0.2"`, while Isaac Lab training has no observation noise in the angular velocity. This created a mismatch where the policy expected clean signals but received noisy ones. **Solution:** Bypass the gyro sensor entirely and read angular velocity directly from `qvel[3:6]`, which provides clean body-frame angular velocity.

Bug 2: MuJoCo Force Limits Not Disabled. The XML specifies `ctrlrange="±23"` and `actuatorfrange="±23"` on joints (and `±5` on wheels), but Isaac Lab training allows 60 N·m (joints) and 36 N·m (wheels). Simply changing the range values is insufficient—the `ctrllimited`, `forcelimited`, and `actfrclimited` flags must be disabled to prevent silent torque clipping. **Solution:** Override flags at initialization:

```
model.actuator_ctrllimited[:] = False
model.actuator_forcelimited[:] = False
```

Bug 3: Wrong Initialization Height. The deployment script initialized the robot at height 0.2562m (base position from XML), but Isaac Lab trains at 0.5562m (defined in `flamingo_rev03_1_1.py:51`). This 0.30m difference (54% error) places the robot in a completely different dynamic regime—the policy’s height control was trained for a tall, extended stance, not a low, compressed stance. **Solution:** Initialize base at $z = 0.5562$ m to match training.

Bug 4: Wheel Axis Inverted Between Simulators. Isaac Lab USD model has **opposite wheel axes** from MuJoCo XML. In MuJoCo: positive wheel velocity = forward motion. In Isaac Lab: positive wheel velocity = backward motion. This axis inversion must be corrected in *both* observations and actions. **Solution:** Negate wheel velocities when reading observations, and negate wheel action targets before applying control:

```
wheel_vel_obs = -qvel[wheel_indices] * 0.15 # Negate for observation
wheel_action_target = -action[6:8] * 20.0 # Negate for control
```

Bug 5: Gear Ratio in PD Control. The leg joints use `GearDelayedPDActuator` where the action target is in *joint space*, but PD comparison happens in *motor space* (scaled by gear). The correct torque formula is:

$$\tau_{\text{motor}} = [K_p(a - q \cdot g) + K_d(0 - \dot{q} \cdot g)] \cdot g \cdot \gamma$$

where $g = -1.5$, γ is the random actuator gain scaling from DR. The final joint torque is $\tau = \tau_{\text{motor}} \cdot g$. Missing any sign or scaling factor breaks leg control.

Bug 6: Isaac Lab Explicit PD Every Substep. Isaac Lab’s `DelayedPDActuator` uses **explicit PD** (not PhysX constraints) and recomputes torques *every physics substep* inside the decimation loop. The action target remains constant for the entire policy step (50 Hz), but PD torques are recalculated 4 times per policy step using the current joint state. MuJoCo deployment must replicate this: apply the same action target for 4 substeps, recomputing PD torques each substep.

Bug 7: CO-RL StateHandler Reset Behavior. On environment reset, the CO-RL framework fills all 3 frames of the history buffer with the *same* initial observation (not zeros). This ensures the policy never sees uninitialized data. Early transfer attempts that zero-initialized the history buffer produced degenerate actions for the first 2 steps (policy seeing 2 zero frames + 1 real frame).

Bug 8: Geometry Mismatch (CRITICAL). All MuJoCo XML files (including `flamingo_torque.xml`, `flamingo_pos_vel.xml`, `flamingo_tester.xml`) have standing height 0.2991m (`base_z = 0.2461m + wheel_radius = 0.053m`), while Isaac Lab training uses init height 0.5562m—a difference of **0.256m (46% error)**. This is not a deployment bug but a

model mismatch: the MuJoCo XML was authored separately from the USD, and they represent different robot geometries. The existing policy is fundamentally incompatible with the MuJoCo model. **Solutions**: (1) Export MuJoCo XML directly from USD at 1:1 scale, OR (2) Retrain policy at 0.30 m height.

4.4.3 Experiment Setup

The transfer was implemented in `transfer_flamingo_sim2sim.py` with all 8 bug fixes applied:

1. Correct 28-dim observation structure matching `env.yaml`
2. Proper PD control with joint-specific gains
3. Joint axis inversion handling (no manual negation needed for XML axes)
4. Gear ratio (-1.5) for leg joint observations and control
5. 3-frame observation stacking (newest first)
6. Gyro sensor bypassed (use `qvel[3:6]` directly)
7. Wheel axis inversion (negate wheel velocities and actions)
8. Correct initialization height (0.5562 m)

Test protocol: Deploy with forward velocity command ($v_x = 0.5$ m/s), rotation command ($\omega_z = 1.0$ rad/s), and standing ($v_x = 0$) for 30 seconds each. The policy used is `Flamingo_Flat_Stand_Drive` (reward 2000+, 5K training iterations, [512, 256, 128] network).

4.4.4 Results

Table 62: Flamingo Transfer Results

Metric	Isaac Lab	MuJoCo	Observation
Balance maintained	Yes	Partial	Maintains for 4+ seconds
Target height (0.45 m)	Achieved	Not achieved	Stands at ~ 0.23 – 0.24 m
Fwd. vel. tracking	Good	Poor	0.2 – 0.3 m/s, ± 0.1 oscillation
Rotation response	Good	Very poor	< 0.1 rad/s ($> 90\%$ tracking error)
Action outputs	Reasonable	Reasonable	Bounded, non-degenerate

Table 63: Flamingo Quantitative Transfer Comparison

Metric	Isaac Lab	MuJoCo
Standing height	0.358 m	0.23–0.24 m
Forward velocity (cmd: 0.5 m/s)	~ 0.48 m/s	0.2–0.3 m/s
Velocity tracking error	$< 5\%$	40–60%
Yaw rate (cmd: 1.0 rad/s)	~ 0.9 rad/s	< 0.1 rad/s
Yaw tracking error	$< 10\%$	$> 90\%$
Training reward	$\sim 2000+$	N/A
Episode length	400 steps (20 s)	Variable (balance-dependent)

4.4.5 Analysis

Dynamic Instability—Supported. The Flamingo transfer is indeed the most fragile of the three robots. The robot maintains basic balance (does not fall over immediately) but fails to achieve meaningful locomotion. Small observation mismatches that Go2 could tolerate (and still walk with) cause the Flamingo to oscillate around its balance point without translating.

Hybrid Control—Supported. The velocity-controlled wheels interact differently with MuJoCo’s physics compared to Isaac Lab:

- Isaac Lab’s **DelayedPDActuator** introduces stochastic delays that the policy is trained to expect. Direct PD in MuJoCo responds instantaneously, creating a timing mismatch in the critical balance-control loop.
- The position-controlled joints and velocity-controlled wheels must coordinate precisely: the legs set body height and balance angle, while the wheels execute translation and rotation. Any desynchronization between these two modes breaks the coordination.

Height Target—Strongly supported. The robot stands at ~ 0.23 – 0.24 m, roughly half the target height of 0.45 m. This height deficit is consistent across all commands and persists throughout the episode. Potential causes include:

- **Gear ratio imprecision:** The -1.5 gear ratio transforms PD gains by $g^2 = 2.25$. Any error in the gear transformation changes the effective leg stiffness and equilibrium height.
- **Different robot model:** Training used USD model `flamingo_rev03_1_1`; MuJoCo uses a separately-authored XML model. Inertial properties, joint limits, and collision geometry may differ.
- **Missing actuator delay:** The policy expects delayed control responses. Instantaneous response may cause it to over-correct, leading to a lower, more conservative stance.

Remaining Issues and Suggested Improvements. Despite fixing all 8 identified bugs, the Flamingo transfer still achieves only partial success (basic balance without meaningful locomotion). The following improvements are prioritized based on their expected impact:

1. **Geometry mismatch resolution (CRITICAL):** The 46% height difference between training (0.5562 m) and deployment model geometry (0.2991 m standing height) is the most fundamental issue. Two solutions:
 - **Preferred:** Export MuJoCo XML directly from Isaac Lab’s USD model using USD-to-MJCF converter at 1:1 scale. This ensures perfect geometry match (mass, COM, inertia, joint limits, collision).
 - **Alternative:** Retrain policy at 0.30 m target height using modified environment config. This is faster but doesn’t address potential inertial property mismatches.
2. **Actuator delay simulation:** Isaac Lab’s **DelayedPDActuator** introduces random 0–4 step delays (uniform distribution, mean = 2 steps). MuJoCo has zero delay. Implement a ring buffer to delay action application by 2 steps:

```
action_buffer = deque(maxlen=3)  # Store 3 most recent actions
action_buffer.append(current_action)
delayed_action = action_buffer[0]  # Use action from 2 steps ago
```

This addresses the timing mismatch in the balance-control loop where the policy expects delayed responses.

3. **Position-velocity actuators:** MuJoCo provides built-in **position** and **velocity** actuator types that may more closely match Isaac Lab’s internal actuator model than manual PD torque computation. The `flamingo_pos_vel.xml` model uses these actuator types. Test transfer with this model and compare against `flamingo_torque.xml`.
4. **PD gain scaling factor:** Similar to Go2’s $6\times$ gain scaling needed to bridge PhysX implicit PD to MuJoCo explicit PD, Flamingo may benefit from a `pd_scale` parameter. The current implementation uses training gains directly; test with `pd_scale` $\in [1, 2, 3, 4, 5, 6]$ to find optimal value. Higher gains may improve height tracking and velocity response.

5. **Gear-transformed PD verification:** The leg joint PD with gear ratio $g = -1.5$ produces effective gains:

$$K_{p,\text{eff}} = g^2 K_p = 2.25 \times 120 = 270 \text{ N}\cdot\text{m}/\text{rad},$$

$$K_{d,\text{eff}} = g^2 K_d = 2.25 \times 1.5 = 3.375 \text{ N}\cdot\text{m}\cdot\text{s}/\text{rad}$$

Validate these effective gains against MuJoCo’s actual joint response with step-response tests (apply constant position target, measure settling time and overshoot).

6. **Observation noise matching:** Training includes observation noise (angular velocity ± 0.15 rad/s, Euler angles ± 0.125 rad), but deployment uses clean sensor data. Add Gaussian noise to match training distribution:

```
ang_vel_obs += np.random.normal(0, 0.15, size=3)
projected_gravity += np.random.normal(0, 0.125, size=3)
```

This prevents the policy from over-reacting to precise observations it was trained to filter.

7. **Domain randomization during deployment:** The policy was trained with actuator gain scaling $\times [0.7, 1.3]$. At deployment, try sweeping this range to find which gain scaling factor produces best tracking. This can compensate for unmodeled dynamics differences without retraining.
8. **Reference model comparison:** The existing `sim2sim_onnx` reference implementation uses different settings (`frame_skip=2`, no gear ratio in PD, no model overrides, `damping=5.0`). Analyze this reference to identify any additional transfer tricks not yet applied to the current implementation.

Priority ranking: (1) Geometry correction is prerequisite for meaningful improvement. (2) Actuator delay addresses the most significant actuator dynamics gap. (3) PD gain scaling and noise matching are lower-cost experiments that may yield incremental improvements. (4) All other items are for fine-tuning after the primary issues are resolved.

4.4.6 Conclusion (Part 3)

The Flamingo two-wheel-legged robot represents the apex of transfer complexity among the three platforms studied. The following conclusions characterize the Flamingo transfer challenge:

1. **Transfer complexity scales super-linearly with mechanical features.** Flamingo required identifying and fixing 8 compounding bugs (vs. 3 for Go2, 8 for Cassie), spanning observation construction, actuator dynamics, geometry mismatch, and control architecture. Each additional mechanical feature (gear ratios, axis inversions, hybrid control, actuator delay, temporal stacking) introduces not just one new potential mismatch, but *interactions* with existing features.
2. **The Flamingo transfer is the most challenging of the three robots.** After all bug fixes, it achieves basic balance (4+ seconds without falling) but fails to produce meaningful locomotion. Standing height is 0.23–0.24 m vs. 0.358 m target (33% deficit). Forward velocity tracking error is 40–60%, and yaw rate tracking error exceeds 90%. This is the poorest transfer quality of any robot in the study.
3. **Dynamic instability amplifies every mismatch.** Small observation or control errors that Go2 tolerates (and still achieves 100% task success with) cause Flamingo to oscillate around its balance point without translating. The balance-control loop has no stability margin: any timing mismatch, gain error, or observation noise immediately degrades locomotion performance. This validates Hypothesis 16.
4. **Hybrid position-velocity control creates additional coupling.** The position-controlled leg joints and velocity-controlled wheel joints must coordinate precisely: legs set body height and balance angle, while wheels execute translation and rotation. Any desynchronization between these two modes (e.g., due to actuator delay mismatch or PD

gain error) breaks the coordination. This is distinct from Go2 and Cassie, which use pure position control. Hypothesis 17 is supported.

5. **Height deficit indicates compound actuator mismatch.** The robot stands at $\sim 50\%$ of target height (0.24 m vs. 0.45 m), consistent across all commands and episodes. This is attributed to: (a) geometry mismatch between USD training model and MuJoCo XML (46% standing height difference), (b) missing actuator delay (policy expects 0–4 step delays, gets zero), and (c) gear ratio imprecision in PD gain transformation. Hypothesis 18 is strongly supported.
6. **Geometry mismatch is a fundamental blocker.** Unlike observation bugs (which can be fixed) or actuator dynamics gaps (which can be bridged with gain scaling), a 46% geometry difference between training and deployment models cannot be resolved through deployment-side corrections. The policy’s learned balance strategy is optimized for a tall, extended stance (0.5562 m init height); attempting to use this policy on a robot with 0.30 m standing height is akin to deploying a giraffe policy on a dachshund. This is the most fundamental lesson from Flamingo: *geometry must match between training and deployment*.
7. **The dominant mismatch source is control architecture.** Flamingo’s primary gap is *control mode coupling and actuator delay*, distinct from Go2 (actuator dynamics: implicit vs. explicit PD) and Cassie (kinematic structure: passive joints, pose conventions). This confirms that the optimal transfer strategy is **morphology-dependent**—there is no universal “one-size-fits-all” approach.
8. **Temporal stacking introduces observation construction complexity.** The 3-frame stacking with newest-first ordering, combined with CO-RL’s reset behavior (fill all frames with initial obs, not zeros), creates subtle but critical transfer requirements. Reversing the frame order or zero-initializing the history produces catastrophically wrong actions. This temporal dimension is absent from Go2 and Cassie, making Flamingo’s observation pipeline the most complex.
9. **The debugging path for Flamingo was non-monotonic.** Unlike Go2 (where fixing PD gains + control swap immediately enabled walking), Flamingo required *all 8 bugs* to be fixed before seeing even partial balance. Early debugging attempts that fixed only 3–5 bugs showed no improvement over random actions, masking progress and making it difficult to identify the critical bugs. This multiplicative bug interaction is characteristic of highly coupled systems.
10. **Flamingo transfer success requires model geometry correction.** The path forward is clear: export MuJoCo XML from USD at 1:1 scale (or retrain at 0.30 m height), then re-test with actuator delay simulation and PD gain scaling. Without geometry correction, further tuning is unlikely to produce meaningful locomotion.

Comparison to Go2 and Cassie: Flamingo’s 40–60% velocity tracking error and $>90\%$ yaw tracking error are dramatically worse than Go2’s 39% RMSE increase (which still achieved 100% task completion) and Cassie’s 1.70s standing (which is a partial success given the passive ankle constraints). Flamingo represents a *qualitative failure* rather than quantitative degradation: the robot maintains balance but does not locomote. This suggests that dynamically unstable platforms are fundamentally more sensitive to sim-gap than statically stable platforms, and may require tighter tolerances on all mismatch dimensions.

5 Cross-Robot Comparative Analysis

5.1 Transfer Complexity Scaling

The three robots represent an increasing gradient of transfer difficulty. Table 64 summarizes the key dimensions:

Table 64: Cross-Robot Transfer Complexity Comparison

Dimension	Go2	Cassie	Flamingo
DOF (train $\rightarrow$ deploy)	12 $\rightarrow$ 12	12 $\rightarrow$ 10	8 $\rightarrow$ 8
Joint mapping	Reorder only	Reorder + drop 2	Reorder + inversions
Pose mismatch	Minimal	≤ 0.6 rad	Moderate
Obs. scaling	Yes ($\times 2.0$, $\times 0.25$)	No (raw values)	Partial ($\times 0.15$, $\times 0.25$)
Temporal stacking	No (single frame)	No (single frame)	Yes (3 frames)
Control mode	Position only	Position only	Position + Velocity
Gear ratios	None	None	-1.5 (legs)
PD gains	Uniform	Variable (3 types)	Variable + gear-scaled
Actuator delay	No	No	Yes (0–4 steps)
Axis inversions	No	No	Yes (right side)
Special constraints	None	4-bar linkage	Dynamic instability
Required Δt	0.005 s (200 Hz)	0.0005 s (2 kHz)	0.005 s (200 Hz)
Bugs found	3	8	8 + geometry
Best survival	100%	1.70 s standing (passive)	Balance only

5.2 Dominant Mismatch Sources by Morphology

Each robot has a different *dominant* mismatch source, suggesting that the optimal transfer strategy is morphology-dependent:

Table 65: Dominant Mismatch Source by Robot

Robot	Dominant Mismatch	Evidence
Go2	Actuator dynamics (implicit vs. explicit PD)	Implicit vs. explicit PD $\rightarrow$ $4\times$ oscillation Linear vel. RMSE $2.93\times$ worse in MuJoCo
Cassie	Kinematic structure (passive joints, defaults)	12 $\rightarrow$ 10 joints, 0.6 rad default offset Obs. baseline: 16 $\rightarrow$ 153 steps; passive: 1.70 s
Flamingo	Control mode coupling (pos/vel hybrid, delay)	Height 0.24 m vs. 0.45 m target Balance ok, locomotion fails

5.3 Common Patterns Across Robots

Despite their different morphologies and dominant mismatch sources, several patterns emerged consistently across all three robots:

1. **Observation construction is the critical path.** In all three cases, the most impactful fixes involved the observation pipeline (joint ordering, scaling, frame conventions, default baselines)—not the physics engine parameters. Getting the “inputs” right matters more than matching the “physics.”
2. **Bugs are multiplicative, not additive.** Multiple simultaneous bugs create compound effects that mask individual contributions. Fixing bugs one at a time may show no improvement until a critical mass of fixes is reached (as seen with Go2’s three simultaneous bugs).
3. **The zero-command test is the essential first diagnostic.** If the robot cannot stand still with zero velocity commands, the problem is in the pipeline, not the policy. This test immediately differentiates pipeline bugs from dynamics mismatch.

4. **Quaternion conventions are a perennial source of bugs.** The sign convention in `quat_rotate_inverse`, the ordering of quaternion components ($[w, x, y, z]$ vs. $[x, y, z, w]$), and the frame convention for angular velocity (`qvel[3:6]`) all caused bugs across the project.
5. **Transfer complexity scales super-linearly with mechanical complexity.** Go2 (simplest mechanism) required 3 fixes. Cassie (4-bar linkage) required 8 (including gain scaling and passive ankle issues). Flamingo (gear ratios + axis inversions + hybrid control + actuator delay + temporal stacking) required 8 bug fixes plus a critical geometry mismatch. Each additional mechanical feature introduces not just one new potential mismatch, but *interactions* with existing features. Flamingo’s 8 bugs span observation (wheel axis inversion, gyro noise), control (gear ratio PD, actuator delay), initialization (height mismatch), and model geometry—the broadest bug surface area of all three robots.

5.4 Domain Randomization Effectiveness

The Go2 experiments (Experiment 1A) provide the most controlled evidence on DR effectiveness:

- **Moderate DR is optimal for the sim-to-sim gap:** It acts as a regularizer that improves generalization (RMSE 0.257 vs. Baseline’s 0.273).
- **Aggressive DR can be counterproductive:** Overly wide randomization degrades multi-axis coordination (Diagonal RMSE 0.453 vs. 0.302).
- **DR addresses contact/mass mismatch but not actuator dynamics:** The 39% sim-gap (Experiment 1B) is dominated by implicit vs. explicit PD, which DR does not randomize.

For Cassie and Flamingo, the transfer quality is dominated by pipeline bugs rather than dynamics mismatch, making DR a secondary concern until the pipeline is fully debugged.

5.5 Complementary Transfer Strategies: DR, ActuatorNet, and RMA

Experiments 1A–1D demonstrated that three different approaches address distinct dimensions of the sim-gap. Table 66 summarizes their characteristics:

Table 66: Mismatch Reduction Strategy Comparison (Go2 Flat Terrain)

Metric	DR (Moderate)	ActuatorNet	RMA	Combined
Overall RMSE	0.257	0.306	(pending)	—
Diagonal RMSE	0.302	0.428	(pending)	—
Diagonal ω_z	0.093	0.323	(pending)	—
Contact robustness	Strong	Weak	Potentially	Strong
Actuator fidelity	Weak	Strong	Potentially	Strong
Online adaptation	No	No	Yes	Yes
Height stability	0.357 m	0.306 m	(pending)	—

Key insight: DR, ActuatorNet, and RMA address *orthogonal* dimensions of the sim-gap:

- **DR** (training-time): Robustifies against contact, friction, and mass variation
- **ActuatorNet** (deployment-time controller): Learns the actual actuator dynamics from MuJoCo data
- **RMA** (deployment-time policy): Adapts the policy’s behavior online by estimating dynamics from observation history

Combining all three—Moderate DR during training, ActuatorNet for actuator fidelity, and RMA

for online adaptation—represents the most complete strategy for bridging the sim-gap, though this combination has not yet been experimentally validated.

The ActuatorNet’s effectiveness on the Diagonal scenario ω_z component (-38%) is particularly noteworthy: this was the exact failure mode that distinguished Moderate DR from Aggressive DR in Experiment 1A, and the ActuatorNet addresses it through a fundamentally different mechanism (learned actuator dynamics vs. parameter randomization).

6 Conclusion

This report presented a comprehensive study of sim-to-sim policy transfer across three robot morphologies—Go2 (quadruped), Cassie (biped), and Flamingo (two-wheel-legged)—from Isaac Lab (PhysX) to MuJoCo. The key findings are:

1. **Transfer pipeline correctness dominates over physics matching.** Across all three robots, the most impactful improvements came from fixing observation construction bugs (joint ordering, frame conventions, default baselines), not from tuning physics parameters. A perfectly matched physics engine with a broken observation pipeline produces worse results than a mismatched physics engine with correct observations.
2. **The dominant mismatch source is morphology-dependent.** Go2’s primary gap is actuator dynamics (implicit vs. explicit PD). Cassie’s is kinematic structure (passive joints, pose conventions, constraint physics). Flamingo’s is control architecture (hybrid pos/vel, actuator delay). A universal “one-size-fits-all” transfer strategy is unlikely to be optimal.
3. **Domain randomization has diminishing returns.** For Go2 flat terrain, Moderate DR (RMSE 0.257) outperforms both Baseline (0.273) and Aggressive (0.271). The robustness-precision trade-off manifests specifically in multi-axis scenarios. DR addresses contact/mass mismatch but not actuator dynamics.
4. **ActuatorNet and RMA represent next-generation solutions.** While Domain Randomization hits a ceiling defined by the robustness-precision trade-off, data-driven methods offer a way to break through. ActuatorNet addresses specific hardware dynamics (like implicit/explicit PD gaps) by learning them directly, while RMA enables policies to adapt online to unmodeled variations. Our initial results indicate these methods can solve the specific oscillation and tracking errors that DR leaves behind.
5. **The sim-gap is moderate but structured.** Go2’s 39% RMSE increase is dominated by systematic oscillation, not random noise. This structured nature makes the gap amenable to targeted solutions like ActuatorNet and RMA, which can model or adapt to these systematic differences better than broad randomization.
6. **Mechanical complexity creates super-linear transfer difficulty.** Each additional mechanical feature (gear ratios, 4-bar linkages, hybrid control, axis inversions) introduces not just one new mismatch, but interactions with existing features. Go2 required 3 bug fixes, Cassie required 8, and Flamingo required 8 bug fixes plus a critical geometry mismatch. Flamingo’s bugs span the broadest surface area: observation (wheel axis inversion, gyro noise), control (gear ratio PD, actuator delay), initialization (height mismatch), and model geometry.
7. **Structural mismatches require architectural solutions.** For Cassie, the passive ankle mismatch (12 trained $\rightarrow$ 10 deployable) cannot be resolved by DR or actuator networks alone—it requires *training with the correct action space* (passive ankle variant with 10 actions and 46 observations). For Flamingo, the 46% geometry mismatch between training model (0.5562 m init height) and deployment model (0.2991 m standing height) cannot be resolved through deployment-side corrections—it requires exporting the MuJoCo XML from USD at 1:1 scale or retraining at the correct height. This principle generalizes: when the training and deployment simulators have different kinematic structures or fundamental

geometry, the models must be matched at the source.

8. **Dynamically unstable platforms are qualitatively more sensitive.** Flamingo’s 40–60% velocity tracking error and >90% yaw tracking error represent a *qualitative failure* (balance without locomotion), unlike Go2’s 39% RMSE increase (which still achieved 100% task completion) and Cassie’s 1.70s standing (partial success). Small observation or control mismatches that statically stable platforms tolerate cause dynamically unstable platforms to oscillate without locomoting. The balance-control loop has no stability margin, amplifying every mismatch dimension. This suggests that two-wheel-legged robots and other dynamically unstable morphologies require tighter tolerances on all transfer pipeline components.
9. **Geometry matching is prerequisite for meaningful transfer.** Flamingo demonstrates that no amount of observation bug fixes, actuator tuning, or PD gain scaling can compensate for a fundamental geometry mismatch. The policy’s learned balance strategy is optimized for a specific robot geometry; attempting to deploy on a robot with 46% different standing height is fundamentally incompatible. This is the most critical lesson from Flamingo: verify geometry match *before* debugging any other transfer issues.

6.1 Future Directions and Methodological Improvements

Based on the experimental findings, three advanced methodological improvements are identified to address the remaining transfer gaps. These recommendations focus on refining the data-driven methods implemented in this study (ActuatorNet, RMA) and resolving the fundamental geometric incompatibility in the Flamingo platform.

6.1.1 Refining Data-Driven Actuator Modeling (ActuatorNet)

While the ActuatorNet (Experiment 1C) successfully reduced tracking error by 38% in complex multi-axis scenarios, it caused a 33.7% performance degradation during steady-state forward walking. This suggests the learned model overfitted to the training distribution of random motor babbling and failed to generalize to the specific dynamics of locomotion.

Proposed Improvement: Locomotion-Centric Data Enrichment

1. **Targeted Data Collection:** Instead of relying solely on random frequency sweeps, the training dataset should be augmented with state-action pairs collected from the target simulator (MuJoCo) during sustained locomotion trajectories (trotting, turning, stopping).
2. **Distribution Matching:** The distribution of joint velocities and torques in the training dataset must explicitly cover the regime occupied by the walking policy.
3. **Hybrid Modeling:** A residual approach could be adopted, where the ActuatorNet predicts only the *error* of the ideal DC motor model ($\tau_{\text{net}} = \tau_{\text{ideal}} + \Delta\tau$), ensuring the physical baseline is preserved for standard motions.

6.1.2 Validating and Extending Rapid Motor Adaptation (RMA)

The implementation of RMA (Experiment 1D) established the two-phase training pipeline. However, the current evaluation on flat terrain does not fully exercise the adaptation module’s capability, as the environment dynamics remain relatively constant.

Proposed Improvement: Rough Terrain and Latent Analysis

1. **Rough Terrain Valuation:** The adaptation module should be evaluated on rough terrain (steps, slopes), where correct estimation of ground reaction forces and friction is critical for survival. This is expected to yield a greater performance separation between RMA and static Domain Randomization.

2. **Latent Space Analysis:** A quantitative analysis of the estimated extrinsic vector $\hat{z}_t$ is required to verify correlation with physical ground truth. This ensures the network is effectively performing system identification rather than memorizing trajectory patterns.

6.1.3 Geometric Retargeting for Flamingo

The Flamingo experiments demonstrated that policy learning cannot overcome the 46% standing height mismatch between the training (USD) and deployment (XML) models.

Proposed Improvement: Unified Asset Pipeline

1. **Single Source of Truth:** Separate manual authoring of USD and XML files must be abandoned. A unified asset pipeline should be implemented where the MuJoCo XML is programmatically generated from the Isaac Lab USD file using a USD-to-MJCF converter.
2. **Automated Verification:** The pipeline must include an automated verification step that computes and asserts equality for total mass, center of mass (CoM) location, and principal moments of inertia between the two representations before training begins.

Expected Impact: Eliminating the geometry mismatch is a prerequisite for Flamingo transfer. Once the kinematic structure is matched, the remaining control challenges (hybrid architecture, delays) can be addressed via ActuatorNet or specific delay injection.

References

- [1] Agility Robotics. “Cassie: A Dynamic Bipedal Robot”. In: (2020). URL: <https://www.agilityrobotics.com/robots/cassie>.
- [2] Jemin Hwangbo, Joonho Lee, and Marco Hutter. “Learning Actuator Network for Sim-to-Real Transfer”. In: *Science Robotics* (2019).
- [3] Jemin Hwangbo et al. “Learning agile and dynamic motor skills for legged robots”. In: *Science Robotics* 4.26 (2019).
- [4] Ashish Kumar et al. “RMA: Rapid Motor Adaptation for Legged Robots”. In: *Robotics: Science and Systems*. 2021.
- [5] Zhongyu Li et al. “Reinforcement Learning for Robust Parameterized Locomotion Control of Bipedal Robots”. In: *IEEE International Conference on Robotics and Automation* (2021).
- [6] Viktor Makoviychuk et al. “Isaac Gym: High Performance GPU-Based Physics Simulation For Robot Learning”. In: *arXiv preprint arXiv:2108.10470*. 2021.
- [7] Gabriel B. Margolis and Pulkit Agrawal. “Walk These Ways: Tuning Robot Control for Generalization with Multiplicity of Behavior”. In: *Conference on Robot Learning*. 2022.
- [8] Takahiro Miki et al. “Learning robust perceptive locomotion for quadrupedal robots in the wild”. In: *Science Robotics* 7.62 (2022).
- [9] Mayank Mittal et al. “Orbit: A Unified Simulation Framework for Interactive Robot Learning Environments”. In: *IEEE Robotics and Automation Letters*. 2023.
- [10] Xue Bin Peng et al. “Sim-to-Real Transfer of Robotic Control with Dynamics Randomization”. In: *IEEE International Conference on Robotics and Automation* (2018).
- [11] Nikita Rudin et al. “Learning to Walk in Minutes Using Massively Parallel Deep Reinforcement Learning”. In: *Conference on Robot Learning*. 2022.
- [12] John Schulman et al. “Proximal Policy Optimization Algorithms”. In: *arXiv preprint arXiv:1707.06347*. 2017.
- [13] Jonah Siekmann et al. “Sim-to-Real Transfer of Blind Bipedal Locomotion for Cassie with Deep Reinforcement Learning”. In: *IEEE Robotics and Automation Letters*. 2021.

- [14] Josh Tobin et al. “Domain Randomization for Transferring Deep Neural Networks from Simulation to the Real World”. In: *IEEE/RSJ International Conference on Intelligent Robots and Systems* (2017).
- [15] Emanuel Todorov, Tom Erez, and Yuval Tassa. “MuJoCo: A physics engine for model-based control”. In: *IEEE/RSJ International Conference on Intelligent Robots and Systems* (2012), pp. 5026–5033.
- [16] Unitree Robotics. “Unitree Go2: A Low-Cost Quadrupedal Robot”. In: (2023). URL: <https://www.unitree.com/go2/>.